

COMPSCIX 415.2 Homework 7

Sanatan Das

March 17, 2018

Contents

Code and Documents Git Repository	2
Load packages (prerequisites to run the code in this document)	2
Analysis of Ames Housing dataset and predicting the price	2

Code and Documents Git Repository

All the work can be found in the below Git repository location:

<https://github.com/sanatanonline/compscix-415-2-assignments>

Load packages (prerequisites to run the code in this document)

```
library(tidyverse)
library(broom)
```

Analysis of Ames Housing dataset and predicting the price

Exercise 1

Load the train.csv dataset into R. How many observations and columns are there?

Answer

```
# Read from train.csv file
train <- read_csv("C:/view/opt/apps/git/R/compscix-415-2-assignments/train.csv")
```

```
## Parsed with column specification:
## cols(
##   .default = col_character(),
##   Id = col_integer(),
##   MSSubClass = col_integer(),
##   LotFrontage = col_integer(),
##   LotArea = col_integer(),
##   OverallQual = col_integer(),
##   OverallCond = col_integer(),
##   YearBuilt = col_integer(),
##   YearRemodAdd = col_integer(),
##   MasVnrArea = col_integer(),
##   BsmtFinSF1 = col_integer(),
##   BsmtFinSF2 = col_integer(),
##   BsmtUnfSF = col_integer(),
##   TotalBsmtSF = col_integer(),
##   `1stFlrSF` = col_integer(),
##   `2ndFlrSF` = col_integer(),
##   LowQualFinSF = col_integer(),
##   GrLivArea = col_integer(),
##   BsmtFullBath = col_integer(),
##   BsmtHalfBath = col_integer(),
##   FullBath = col_integer()
##   # ... with 18 more columns
## )

## See spec(...) for full column specifications.

# glimpse train
glimpse(train)
```

```

## Observations: 1,460
## Variables: 81
## $ Id <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 1...
## $ MSSubClass <int> 60, 20, 60, 70, 60, 50, 20, 60, 50, 190, 20, 60,...
## $ MSZoning <chr> "RL", "RL", "RL", "RL", "RL", "RL", "RL", "RL", ...
## $ LotFrontage <int> 65, 80, 68, 60, 84, 85, 75, NA, 51, 50, 70, 85, ...
## $ LotArea <int> 8450, 9600, 11250, 9550, 14260, 14115, 10084, 10...
## $ Street <chr> "Pave", "Pave", "Pave", "Pave", "Pave", "Pave", ...
## $ Alley <chr> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, ...
## $ LotShape <chr> "Reg", "Reg", "IR1", "IR1", "IR1", "IR1", "Reg",...
## $ LandContour <chr> "Lvl", "Lvl", "Lvl", "Lvl", "Lvl", "Lvl", "Lvl",...
## $ Utilities <chr> "AllPub", "AllPub", "AllPub", "AllPub", "AllPub"...
## $ LotConfig <chr> "Inside", "FR2", "Inside", "Corner", "FR2", "Ins...
## $ LandSlope <chr> "Gtl", "Gtl", "Gtl", "Gtl", "Gtl", "Gtl", "Gtl",...
## $ Neighborhood <chr> "CollgCr", "Veenker", "CollgCr", "Crawfor", "NoR...
## $ Condition1 <chr> "Norm", "Feedr", "Norm", "Norm", "Norm", "Norm",...
## $ Condition2 <chr> "Norm", "Norm", "Norm", "Norm", "Norm", "Norm", ...
## $ BldgType <chr> "1Fam", "1Fam", "1Fam", "1Fam", "1Fam", "1Fam", ...
## $ HouseStyle <chr> "2Story", "1Story", "2Story", "2Story", "2Story"...
## $ OverallQual <int> 7, 6, 7, 7, 8, 5, 8, 7, 7, 5, 5, 9, 5, 7, 6, 7, ...
## $ OverallCond <int> 5, 8, 5, 5, 5, 5, 5, 6, 5, 6, 5, 5, 6, 5, 5, 8, ...
## $ YearBuilt <int> 2003, 1976, 2001, 1915, 2000, 1993, 2004, 1973, ...
## $ YearRemodAdd <int> 2003, 1976, 2002, 1970, 2000, 1995, 2005, 1973, ...
## $ RoofStyle <chr> "Gable", "Gable", "Gable", "Gable", "Gable", "Ga...
## $ RoofMatl <chr> "CompShg", "CompShg", "CompShg", "CompShg", "Com...
## $ Exterior1st <chr> "VinylSd", "MetalSd", "VinylSd", "Wd Sdng", "Vin...
## $ Exterior2nd <chr> "VinylSd", "MetalSd", "VinylSd", "Wd Shng", "Vin...
## $ MasVnrType <chr> "BrkFace", "None", "BrkFace", "None", "BrkFace",...
## $ MasVnrArea <int> 196, 0, 162, 0, 350, 0, 186, 240, 0, 0, 0, 286, ...
## $ ExterQual <chr> "Gd", "TA", "Gd", "TA", "Gd", "TA", "Gd", "TA", ...
## $ ExterCond <chr> "TA", "TA", "TA", "TA", "TA", "TA", "TA", "TA", ...
## $ Foundation <chr> "PConc", "CBlock", "PConc", "BrkTil", "PConc", "...
## $ BsmtQual <chr> "Gd", "Gd", "Gd", "TA", "Gd", "Gd", "Ex", "Gd", ...
## $ BsmtCond <chr> "TA", "TA", "TA", "Gd", "TA", "TA", "TA", "TA", ...
## $ BsmtExposure <chr> "No", "Gd", "Mn", "No", "Av", "No", "Av", "Mn", ...
## $ BsmtFinType1 <chr> "GLQ", "ALQ", "GLQ", "ALQ", "GLQ", "GLQ", "GLQ",...
## $ BsmtFinSF1 <int> 706, 978, 486, 216, 655, 732, 1369, 859, 0, 851,...
## $ BsmtFinType2 <chr> "Unf", "Unf", "Unf", "Unf", "Unf", "Unf", "Unf",...
## $ BsmtFinSF2 <int> 0, 0, 0, 0, 0, 0, 0, 32, 0, 0, 0, 0, 0, 0, 0, ...
## $ BsmtUnfSF <int> 150, 284, 434, 540, 490, 64, 317, 216, 952, 140,...
## $ TotalBsmtSF <int> 856, 1262, 920, 756, 1145, 796, 1686, 1107, 952,...
## $ Heating <chr> "GasA", "GasA", "GasA", "GasA", "GasA", "GasA", ...
## $ HeatingQC <chr> "Ex", "Ex", "Ex", "Gd", "Ex", "Ex", "Ex", "Ex", ...
## $ CentralAir <chr> "Y", "Y", "Y", "Y", "Y", "Y", "Y", "Y", "Y", "Y", "Y"...
## $ Electrical <chr> "SBrkr", "SBrkr", "SBrkr", "SBrkr", "SBrkr", "SB...
## $ `1stFlrSF` <int> 856, 1262, 920, 961, 1145, 796, 1694, 1107, 1022...
## $ `2ndFlrSF` <int> 854, 0, 866, 756, 1053, 566, 0, 983, 752, 0, 0, ...
## $ LowQualFinSF <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
## $ GrLivArea <int> 1710, 1262, 1786, 1717, 2198, 1362, 1694, 2090, ...
## $ BsmtFullBath <int> 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, ...
## $ BsmtHalfBath <int> 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
## $ FullBath <int> 2, 2, 2, 1, 2, 1, 2, 2, 2, 1, 1, 3, 1, 2, 1, ...
## $ HalfBath <int> 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, ...
## $ BedroomAbvGr <int> 3, 3, 3, 3, 4, 1, 3, 3, 2, 2, 3, 4, 2, 3, 2, 2, ...

```

```
## $ KitchenAbvGr <int> 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 1, 1, 1, 1, 1, ...
## $ KitchenQual <chr> "Gd", "TA", "Gd", "Gd", "Gd", "TA", "Gd", "TA", ...
## $ TotRmsAbvGrd <int> 8, 6, 6, 7, 9, 5, 7, 7, 8, 5, 5, 11, 4, 7, 5, ...
## $ Functional <chr> "Typ", "Typ", "Typ", "Typ", "Typ", "Typ", "Typ", ...
## $ Fireplaces <int> 0, 1, 1, 1, 1, 0, 1, 2, 2, 2, 0, 2, 0, 1, 1, 0, ...
## $ FireplaceQu <chr> NA, "TA", "TA", "Gd", "TA", NA, "Gd", "TA", "TA"...
## $ GarageType <chr> "Attchd", "Attchd", "Attchd", "Detchd", "Attchd"...
## $ GarageYrBlt <int> 2003, 1976, 2001, 1998, 2000, 1993, 2004, 1973, ...
## $ GarageFinish <chr> "RFn", "RFn", "RFn", "Unf", "RFn", "Unf", "RFn", ...
## $ GarageCars <int> 2, 2, 2, 3, 3, 2, 2, 2, 2, 1, 1, 3, 1, 3, 1, 2, ...
## $ GarageArea <int> 548, 460, 608, 642, 836, 480, 636, 484, 468, 205...
## $ GarageQual <chr> "TA", "TA", "TA", "TA", "TA", "TA", "TA", "TA", ...
## $ GarageCond <chr> "TA", "TA", "TA", "TA", "TA", "TA", "TA", "TA", ...
## $ PavedDrive <chr> "Y", "Y", "Y", "Y", "Y", "Y", "Y", "Y", "Y", "Y", "Y"...
## $ WoodDeckSF <int> 0, 298, 0, 0, 192, 40, 255, 235, 90, 0, 0, 147, ...
## $ OpenPorchSF <int> 61, 0, 42, 35, 84, 30, 57, 204, 0, 4, 0, 21, 0, ...
## $ EnclosedPorch <int> 0, 0, 0, 272, 0, 0, 0, 228, 205, 0, 0, 0, 0, 0, ...
## $ `3SsnPorch` <int> 0, 0, 0, 0, 0, 320, 0, 0, 0, 0, 0, 0, 0, 0, 0...
## $ ScreenPorch <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 176, 0, 0, 0...
## $ PoolArea <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
## $ PoolQC <chr> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, ...
## $ Fence <chr> NA, NA, NA, NA, NA, "MnPrv", NA, NA, NA, NA, NA, ...
## $ MiscFeature <chr> NA, NA, NA, NA, NA, "Shed", NA, "Shed", NA, NA, ...
## $ MiscVal <int> 0, 0, 0, 0, 0, 700, 0, 350, 0, 0, 0, 0, 0, 0, 0, ...
## $ MoSold <int> 2, 5, 9, 2, 12, 10, 8, 11, 4, 1, 2, 7, 9, 8, 5, ...
## $ YrSold <int> 2008, 2007, 2008, 2006, 2008, 2009, 2007, 2009, ...
## $ SaleType <chr> "WD", "WD", "WD", "WD", "WD", "WD", "WD", "WD", ...
## $ SaleCondition <chr> "Normal", "Normal", "Normal", "Abnorml", "Normal...
## $ SalePrice <int> 208500, 181500, 223500, 140000, 250000, 143000, ...
```

So there are 1460 observations with 81 columns (variables)

Exercise 2

Normally at this point you would spend a few days on EDA, but for this homework we will get right to fitting some linear regression models. Our first step is to randomly split the data into train and test datasets. We will use a 70/30 split. There is an R package that will do the split for you, but let's get some more practice with R and do it ourselves by filling in the blanks in the code below.

```
# load packages
library(tidyverse)
library(broom)

# When taking a random sample, it is often useful to set a seed so that
# your work is reproducible. Setting a seed will guarantee that the same
# random sample will be generated every time, so long as you always set the
# same seed beforehand
set.seed(29283)

# This data already has an Id column which we can make use of.
# Let's create our training set using sample_frac. Fill in the blank.
train_set <- train %>% sample_frac(____)

# let's create our testing set using the Id column. Fill in the blanks.
test_set <- train %>% filter(!(____ %in% ____$Id))
```

Answer

Let's fill in the blanks.

```
# When taking a random sample, it is often useful to set a seed so that
# your work is reproducible. Setting a seed will guarantee that the same
# random sample will be generated every time, so long as you always set the
# same seed beforehand
```

```
set.seed(29283)
```

```
# This data already has an Id column which we can make use of.
```

```
# Let's create our training set using sample_frac. Fill in the blank.
```

```
train_set <- train %>% sample_frac(0.7)
```

```
# Print train set
```

```
train_set
```

```
## # A tibble: 1,022 x 81
```

```
##       Id MSSubClass MSZoning LotFrontage LotArea Street Alley LotShape
```

```
##    <int>      <int> <chr>          <int>    <int> <chr>  <chr> <chr>
```

```
##  1     22         45 RM              57     7449 Pave   Grv1  Reg
```

```
##  2    637         30 RM              51     6120 Pave   <NA>  Reg
```

```
##  3    121         80 RL              NA    21453 Pave   <NA>  IR1
```

```
##  4    575         80 RL              70    10500 Pave   <NA>  Reg
```

```
##  5   1423        120 RM              37     4435 Pave   <NA>  Reg
```

```
##  6   1169         70 RL             120    13728 Pave   <NA>  Reg
```

```
##  7   1261         60 RL              NA    24682 Pave   <NA>  IR3
```

```
##  8   1319         20 RL              NA    14781 Pave   <NA>  IR2
```

```
##  9     116        160 FV              34     3230 Pave   Pave  Reg
```

```
## 10   1125         80 RL              NA     9125 Pave   <NA>  IR1
```

```
## # ... with 1,012 more rows, and 73 more variables: LandContour <chr>,
```

```
## # Utilities <chr>, LotConfig <chr>, LandSlope <chr>, Neighborhood <chr>,
```

```
## # Condition1 <chr>, Condition2 <chr>, BldgType <chr>, HouseStyle <chr>,
```

```
## # OverallQual <int>, OverallCond <int>, YearBuilt <int>,
```

```
## # YearRemodAdd <int>, RoofStyle <chr>, RoofMatl <chr>,
```

```
## # Exterior1st <chr>, Exterior2nd <chr>, MasVnrType <chr>,
```

```
## # MasVnrArea <int>, ExterQual <chr>, ExterCond <chr>, Foundation <chr>,
```

```
## # BsmtQual <chr>, BsmtCond <chr>, BsmtExposure <chr>,
```

```
## # BsmtFinType1 <chr>, BsmtFinSF1 <int>, BsmtFinType2 <chr>,
```

```
## # BsmtFinSF2 <int>, BsmtUnfSF <int>, TotalBsmtSF <int>, Heating <chr>,
```

```
## # HeatingQC <chr>, CentralAir <chr>, Electrical <chr>, `1stFlrSF` <int>,
```

```
## # `2ndFlrSF` <int>, LowQualFinSF <int>, GrLivArea <int>,
```

```
## # BsmtFullBath <int>, BsmtHalfBath <int>, FullBath <int>,
```

```
## # HalfBath <int>, BedroomAbvGr <int>, KitchenAbvGr <int>,
```

```
## # KitchenQual <chr>, TotRmsAbvGrd <int>, Functional <chr>,
```

```
## # Fireplaces <int>, FireplaceQu <chr>, GarageType <chr>,
```

```
## # GarageYrBlt <int>, GarageFinish <chr>, GarageCars <int>,
```

```
## # GarageArea <int>, GarageQual <chr>, GarageCond <chr>,
```

```
## # PavedDrive <chr>, WoodDeckSF <int>, OpenPorchSF <int>,
```

```
## # EnclosedPorch <int>, `3SsnPorch` <int>, ScreenPorch <int>,
```

```
## # PoolArea <int>, PoolQC <chr>, Fence <chr>, MiscFeature <chr>,
```

```
## # MiscVal <int>, MoSold <int>, YrSold <int>, SaleType <chr>,
```

```
## # SaleCondition <chr>, SalePrice <int>
```

```
# let's create our testing set using the Id column. Fill in the blanks.
```

```
test_set <- train %>% filter(!(train$Id %in% train_set$Id))
```

```
# Print test set
```

```
test_set
```

```
## # A tibble: 438 x 81
```

```
##       Id MSSubClass MSZoning LotFrontage LotArea Street Alley LotShape
```

```
##      <int>      <int> <chr>      <int>      <int> <chr> <chr> <chr>
## 1      1      60 RL      65      8450 Pave <NA> Reg
## 2      2      20 RL      80      9600 Pave <NA> Reg
## 3      3      60 RL      68     11250 Pave <NA> IR1
## 4      4      70 RL      60      9550 Pave <NA> IR1
## 5     14      20 RL      91     10652 Pave <NA> IR1
## 6     23      20 RL      75      9742 Pave <NA> Reg
## 7     27      20 RL      60      7200 Pave <NA> Reg
## 8     38      20 RL      74      8532 Pave <NA> Reg
## 9     40      90 RL      65      6040 Pave <NA> Reg
## 10    42      20 RL     115     16905 Pave <NA> Reg
## # ... with 428 more rows, and 73 more variables: LandContour <chr>,
## #   Utilities <chr>, LotConfig <chr>, LandSlope <chr>, Neighborhood <chr>,
## #   Condition1 <chr>, Condition2 <chr>, BldgType <chr>, HouseStyle <chr>,
## #   OverallQual <int>, OverallCond <int>, YearBuilt <int>,
## #   YearRemodAdd <int>, RoofStyle <chr>, RoofMatl <chr>,
## #   Exterior1st <chr>, Exterior2nd <chr>, MasVnrType <chr>,
## #   MasVnrArea <int>, ExterQual <chr>, ExterCond <chr>, Foundation <chr>,
## #   BsmtQual <chr>, BsmtCond <chr>, BsmtExposure <chr>,
## #   BsmtFinType1 <chr>, BsmtFinSF1 <int>, BsmtFinType2 <chr>,
## #   BsmtFinSF2 <int>, BsmtUnfSF <int>, TotalBsmtSF <int>, Heating <chr>,
## #   HeatingQC <chr>, CentralAir <chr>, Electrical <chr>, `1stFlrSF` <int>,
## #   `2ndFlrSF` <int>, LowQualFinSF <int>, GrLivArea <int>,
## #   BsmtFullBath <int>, BsmtHalfBath <int>, FullBath <int>,
## #   HalfBath <int>, BedroomAbvGr <int>, KitchenAbvGr <int>,
## #   KitchenQual <chr>, TotRmsAbvGrd <int>, Functional <chr>,
## #   Fireplaces <int>, FireplaceQu <chr>, GarageType <chr>,
## #   GarageYrBlt <int>, GarageFinish <chr>, GarageCars <int>,
## #   GarageArea <int>, GarageQual <chr>, GarageCond <chr>,
## #   PavedDrive <chr>, WoodDeckSF <int>, OpenPorchSF <int>,
## #   EnclosedPorch <int>, `3SsnPorch` <int>, ScreenPorch <int>,
## #   PoolArea <int>, PoolQC <chr>, Fence <chr>, MiscFeature <chr>,
## #   MiscVal <int>, MoSold <int>, YrSold <int>, SaleType <chr>,
## #   SaleCondition <chr>, SalePrice <int>
```

Now, we have separated our train data set and test data set.

Exercise 3

Our target is called SalePrice. First, we can fit a simple regression model consisting of only the intercept (the average of SalePrice). Fit the model and then use the broom package to

- take a look at the coefficient,
- compare the coefficient to the average value of SalePrice, and
- take a look at the R-squared.

Use the code below and fill in the blanks.

```
# Fit a model with intercept only
mod_0 <- lm(SalePrice ~ 1, data = _____)
# Double-check that the average SalePrice is equal to our model's coefficient
mean(train_set$SalePrice)
tidy(____)
# Check the R-squared
glance(____)
```

Answer

Let's fill in the blanks.

```
# Fit a model with intercept only
mod_0 <- lm(SalePrice ~ 1, data = train_set)
# Double-check that the average SalePrice is equal to our model's coefficient
mean(train_set$SalePrice)
```

```
## [1] 182176
```

```
tidy(mod_0)
```

```
##           term estimate std.error statistic p.value
## 1 (Intercept)   182176   2492.072    73.10222      0
```

```
# Check the R-squared
```

```
glance(mod_0)
```

```
##   r.squared adj.r.squared   sigma statistic p.value df   logLik   AIC
## 1         0             0 79668.37         NA      NA  1 -12983.57 25971.13
##           BIC      deviance df.residual
## 1 25980.99 6.480338e+12          1021
```

Exercise 4

Now fit a linear regression model using GrLivArea, OverallQual, and Neighborhood as the features. Don't forget to look at data_description.txt to understand what these variables mean. Ask yourself these questions before fitting the model:

- What kind of relationship will these features have with our target?
- Can the relationship be estimated linearly?
- Are these good features, given the problem we are trying to solve?

After fitting the model, output the coefficients and the R-squared using the broom package.


Answer these questions:

- How would you interpret the coefficients on GrLivArea and OverallQual?
- How would you interpret the coefficient on NeighborhoodBrkSide?
- Are the features significant?
- Are the features practically significant?
- Is the model a good fit (to the training set)?

Answer


Let's plot the graph to see the relationship between SalePrice and GrLivArea, OverallQual, and Neighborhood.

```
ggplot(train_set, aes(x = GrLivArea, y = SalePrice)) +
  geom_point(color = "blue") +
  geom_smooth(method = "lm", se = FALSE, color = "red") +
  theme_bw()
```




The relationship between SalePrice and GrLivArea is linear. It can be estimated linearly. It is an important feature to estimate as it impacts the SalePrice.

```
ggplot(train_set, aes(x = OverallQual, y = SalePrice)) +  
  geom_point(color = "blue") +  
  geom_smooth(method = "lm", se = FALSE, color = "red") +  
  theme_bw()
```

The relationship between SalePrice and OverallQual is linear. It can be estimated linearly. It is an important feature to estimate as it impacts the SalePrice.

```
train_set %>%
ggplot() +
  geom_bar(aes(x = Neighborhood, y = SalePrice), stat = 'identity') +
  theme_bw()
```



There is a relationship between SalePrice and Neighborhood. But the Neighborhood is categorical, so no linear relationship with SalePrice. It is an important feature to estimate as it impacts the SalePrice. We will factor it and add to our regression model.

Let's create the models and see the values.

```
# Fit a model for GrLivArea
```

```
lm_1 <- lm(SalePrice ~ GrLivArea, data = train_set)
```

```
mean(train_set$SalePrice)
```

```
## [1] 182176
```

```
tidy(lm_1)
```

```
##      term estimate  std.error statistic    p.value
## 1 (Intercept) 7518.567 5334.586903    1.40940 1.590217e-01
## 2   GrLivArea  115.833   3.354991   34.52558 1.189915e-173
```

```
glance(lm_1)
```

```
##   r.squared adj.r.squared  sigma statistic    p.value df logLik  AIC
## 1 0.5388821    0.5384301 54125.85 1192.016 1.189915e-173  2 -12588 25182
##      BIC  deviance df.residual
## 1 25196.79 2.9882e+12        1020
```

```
# Fit a model for OverallQual
```

```
lm_2 <- lm(SalePrice ~ OverallQual, data = train_set)
```

```
mean(train_set$SalePrice)
```

```
## [1] 182176
```

```
tidy(lm_2)
```

```
##           term estimate std.error statistic      p.value
## 1 (Intercept) -105872.5  6893.466 -15.35839  4.801210e-48
## 2 OverallQual   47192.3  1102.649  42.79902  5.918164e-230
```

```
glance(lm_2)
```

```
##   r.squared adj.r.squared   sigma statistic      p.value df    logLik
## 1 0.6423257    0.641975 47669.72  1831.756 5.918164e-230  2 -12458.19
##           AIC      BIC    deviance df.residual
## 1 24922.38 24937.17 2.317851e+12         1020
```

```
# Fit a model for Neighborhood
```

```
train_set <- train_set %>% mutate(Neighborhood_fct = factor(Neighborhood, ordered = FALSE))
lm_3 <- lm(SalePrice ~ Neighborhood_fct, data = train_set)
mean(train_set$SalePrice)
```

```
## [1] 182176
```

```
tidy(lm_3)
```

```
##           term estimate std.error statistic      p.value
## 1 (Intercept) 200308.462  14880.25  13.4613631 4.487217e-38
## 2 Neighborhood_fctBlueste -76308.462  55676.80 -1.3705612 1.708203e-01
## 3 Neighborhood_fctBrDale -92908.462  21979.59 -4.2270339 2.585277e-05
## 4 Neighborhood_fctBrkSide -76713.249  16812.68 -4.5628209 5.675519e-06
## 5 Neighborhood_fctClearCr  19853.672  20330.29  0.9765561 3.290259e-01
## 6 Neighborhood_fctCollgCr -4491.700  15774.54 -0.2847437 7.758997e-01
## 7 Neighborhood_fctCrawfor  4905.221  17077.14  0.2872390 7.739890e-01
## 8 Neighborhood_fctEdwards -69405.077  16300.50 -4.2578500 2.259005e-05
## 9 Neighborhood_fctGilbert -9247.988  16490.05 -0.5608225 5.750446e-01
## 10 Neighborhood_fctIDOTRR -94072.098  18768.65 -5.0121942 6.365074e-07
## 11 Neighborhood_fctMeadowV -103085.385  21043.85 -4.8985984 1.125722e-06
## 12 Neighborhood_fctMitchel -44658.462  17727.84 -2.5191151 1.192047e-02
## 13 Neighborhood_fctNAMES -54161.231  15455.33 -3.5043723 4.780249e-04
## 14 Neighborhood_fctNoRidge 135254.147  18616.48  7.2652908 7.499696e-13
## 15 Neighborhood_fctNPkVill -59058.462  25152.21 -2.3480422 1.906702e-02
## 16 Neighborhood_fctNridgHt 118035.386  16438.06  7.1806164 1.355530e-12
## 17 Neighborhood_fctNWAmes -12271.923  16636.63 -0.7376449 4.609038e-01
## 18 Neighborhood_fctOldTown -72380.651  16134.44 -4.4860949 8.097746e-06
## 19 Neighborhood_fctSawyer -63192.622  16703.04 -3.7833005 1.639801e-04
## 20 Neighborhood_fctSawyerW -20425.787  16981.27 -1.2028425 2.293228e-01
## 21 Neighborhood_fctSomerst  34565.772  16413.31  2.1059604 3.545731e-02
## 22 Neighborhood_fctStoneBr 118543.907  19311.16  6.1386206 1.199918e-09
## 23 Neighborhood_fctSWISU -52975.087  20033.15 -2.6443711 8.312879e-03
## 24 Neighborhood_fctTimber  52345.885  18224.51  2.8722792 4.161598e-03
## 25 Neighborhood_fctVeenker  63620.110  25152.21  2.5294039 1.157881e-02
```

```
glance(lm_3)
```

```
##   r.squared adj.r.squared   sigma statistic      p.value df    logLik
## 1 0.5571452    0.5464847 53651.51  52.26259 1.011167e-157 25 -12567.35
##           AIC      BIC    deviance df.residual
## 1 25186.7 25314.87 2.869849e+12         997
```

```
# Fit a model with all three variables
lm_4 <- lm(SalePrice ~ GrLivArea + OverallQual + Neighborhood_fct, data = train_set)
mean(train_set$SalePrice)
```

```
## [1] 182176
```

```
tidy(lm_4)
```

##	term	estimate	std.error	statistic
## 1	(Intercept)	-45017.87483	12933.341808	-3.4807612
## 2	GrLivArea	62.77735	3.006033	20.8837885
## 3	OverallQual	21692.23178	1353.714104	16.0242342
## 4	Neighborhood_fctBlueste	-38288.88063	36531.907177	-1.0480942
## 5	Neighborhood_fctBrkDale	-43314.05372	14524.693991	-2.9820975
## 6	Neighborhood_fctBrkSide	-14064.37052	11318.850018	-1.2425618
## 7	Neighborhood_fctClearCr	27839.00662	13561.346871	2.0528202
## 8	Neighborhood_fctCollgCr	4297.67432	10372.304467	0.4143413
## 9	Neighborhood_fctCrawfor	7423.05573	11371.511784	0.6527765
## 10	Neighborhood_fctEdwards	-15284.11495	10994.287187	-1.3901870
## 11	Neighborhood_fctGilbert	-8357.55930	10894.173472	-0.7671586
## 12	Neighborhood_fctIDOTRR	-32689.43085	12603.712743	-2.5936350
## 13	Neighborhood_fctMeadowV	-14446.06504	14190.148622	-1.0180348
## 14	Neighborhood_fctMitchel	1922.31487	11788.608170	0.1630655
## 15	Neighborhood_fctNAMES	-7719.67883	10375.956174	-0.7439969
## 16	Neighborhood_fctNoRidge	47685.16790	12567.432633	3.7943444
## 17	Neighborhood_fctNPkVill	-20240.71145	16548.664867	-1.2231024
## 18	Neighborhood_fctNridgHt	63872.80848	10880.456671	5.8704161
## 19	Neighborhood_fctNWAmes	-12279.33299	11047.502893	-1.1115030
## 20	Neighborhood_fctOldTown	-36107.07577	10849.170903	-3.3280954
## 21	Neighborhood_fctSawyer	-4121.92502	11252.369778	-0.3663162
## 22	Neighborhood_fctSawyerW	-5391.97074	11230.758221	-0.4801075
## 23	Neighborhood_fctSomerst	18700.96725	10772.212794	1.7360377
## 24	Neighborhood_fctStoneBr	65712.45881	12745.312907	5.1558137
## 25	Neighborhood_fctSWISU	-45451.86707	13564.792586	-3.3507233
## 26	Neighborhood_fctTimber	27925.08619	11985.325857	2.3299397
## 27	Neighborhood_fctVeenker	54913.12768	16521.075497	3.3238228
##	p.value			
## 1		5.216927e-04		
## 2		1.337222e-80		
## 3		1.389020e-51		
## 4		2.948497e-01		
## 5		2.932566e-03		
## 6		2.143221e-01		
## 7		4.035110e-02		
## 8		6.787135e-01		
## 9		5.140512e-01		
## 10		1.647830e-01		
## 11		4.431692e-01		
## 12		9.636216e-03		
## 13		3.089089e-01		
## 14		8.705000e-01		
## 15		4.570540e-01		
## 16		1.569690e-04		
## 17		2.215806e-01		

```
## 18 5.917964e-09
## 19 2.666204e-01
## 20 9.064637e-04
## 21 7.142070e-01
## 22 6.312565e-01
## 23 8.286672e-02
## 24 3.045915e-07
## 25 8.363074e-04
## 26 2.000859e-02
## 27 9.203087e-04
```

```
glance(lm_4)
```

```
##   r.squared adj.r.squared   sigma statistic p.value df   logLik   AIC
## 1 0.8099927    0.8050277 35178.1  163.1401      0 27 -12134.95 24325.91
##           BIC      deviance df.residual
## 1 24463.93 1.231311e+12          995
```

How would you interpret the coefficients on GrLivArea and OverallQual?

Answer

How would you interpret the coefficient on NeighborhoodBrkSide?

Answer

Are the features significant?

Answer

Are the features practically significant?

Answer

Is the model a good fit (to the training set)?

Answer

Exercise 5

Evaluate the model on test_set using the root mean squared error (RMSE). Use the predict function to get the model predictions for the testing set.

Hint: use the sqrt() and mean() functions:

```
test_predictions <- predict(NAME_OF_YOUR_MODEL_HERE, newdata = test_set)
rmse <- sqrt(mean((___ - ___)^2))
```

Answer

Let's predict on the test data and evaluate the model.

```
test_set <- test_set %>% mutate(Neighborhood_fct = factor(Neighborhood, ordered = FALSE))
test_predictions <- predict(lm_4, newdata = test_set)
rmse <- sqrt(mean((test_set$SalePrice - mean(test_set$SalePrice))^2))
rmse
```

```
## [1] 78835.85
```

Exercise 7

One downside of the linear model is that it is sensitive to unusual values because the distance incorporates a squared term. Fit a linear model to the simulated data below, and visualise


the results. Rerun a few times to generate different simulated datasets. What do you notice about the model?

```
sim1a <- tibble(  
  x = rep(1:10, each = 3),  
  y = x * 1.5 + 6 + rt(length(x), df = 2)  
)
```

Answer

Lets create a model and run on the simulated data and visualize it.


```
sim1a <- tibble(  
  x = rep(1:10, each = 3),  
  y = x * 1.5 + 6 + rt(length(x), df = 2)  
)  
  
mod_5 <- lm(y~x, data = sim1a)  
ggplot(sim1a,aes(x,y))+  
  geom_point(size = 2, color = "blue")+  
  geom_abline(intercept = mod_5$coefficients[1],slope = mod_5$coefficients[2])
```



Now, let's run for few times.


```
sim2a <- tibble(  
  x = rep(1:10, each = 3),  
  y = x * 1.5 + 6 + rt(length(x), df = 2)  
)
```

```
mod_5 <- lm(y~x, data = sim1a)
ggplot(sim2a,aes(x,y))+
  geom_point(size = 2, color = "blue")+
  geom_abline(intercept = mod_5$coefficients[1],slope = mod_5$coefficients[2])
```



```
sim3a <- tibble(
  x = rep(1:10, each = 3),
  y = x * 1.5 + 6 + rt(length(x), df = 2)
)

mod_5 <- lm(y~x, data = sim1a)
ggplot(sim3a,aes(x,y))+
  geom_point(size = 2, color = "blue")+
  geom_abline(intercept = mod_5$coefficients[1],slope = mod_5$coefficients[2])
```



Conclusion: Sometimes, one single abnormal value forces the fitted line deviate from the “intuitively” best lines.

End of Homework 7