

COMPSCIX 415.2 Homework 7

Sanatan Das

March 17, 2018

Contents

Code and Documents Git Repository	2
Load packages (prerequisites to run the code in this document)	2
Analysis of Ames Housing dataset and predicting the price	2
Load the data	2
Split the Data to training set and test set	4
The intercept: <i>SalePrice</i>	6
EDA on <i>GrLivArea</i> , <i>OverallQual</i> , and <i>Neighborhood</i>	7
Evaluate the model	17
Linear Model : downside	17

Code and Documents Git Repository

All the work can be found in the below Git repository location:

<https://github.com/sanatanonline/compscix-415-2-assignments>

Load packages (prerequisites to run the code in this document)

```
library(tidyverse)
library(broom)
```

Analysis of Ames Housing dataset and predicting the price

Load the data

Exercise 1

Load the train.csv dataset into R. How many observations and columns are there?

Answer

```
# Read from train.csv file
train <- read_csv("C:/view/opt/apps/git/R/compscix-415-2-assignments/train.csv")
```

```
## Parsed with column specification:
## cols(
##   .default = col_character(),
##   Id = col_integer(),
##   MSSubClass = col_integer(),
##   LotFrontage = col_integer(),
##   LotArea = col_integer(),
##   OverallQual = col_integer(),
##   OverallCond = col_integer(),
##   YearBuilt = col_integer(),
##   YearRemodAdd = col_integer(),
##   MasVnrArea = col_integer(),
##   BsmtFinSF1 = col_integer(),
##   BsmtFinSF2 = col_integer(),
##   BsmtUnfSF = col_integer(),
##   TotalBsmtSF = col_integer(),
##   `1stFlrSF` = col_integer(),
##   `2ndFlrSF` = col_integer(),
##   LowQualFinSF = col_integer(),
##   GrLivArea = col_integer(),
##   BsmtFullBath = col_integer(),
##   BsmtHalfBath = col_integer(),
##   FullBath = col_integer()
##   # ... with 18 more columns
## )

## See spec(...) for full column specifications.
```

```
# glimpse train
glimpse(train)
```

```
## Observations: 1,460
## Variables: 81
## $ Id <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 1...
## $ MSSubClass <int> 60, 20, 60, 70, 60, 50, 20, 60, 50, 190, 20, 60,...
## $ MSZoning <chr> "RL", "RL", "RL", "RL", "RL", "RL", "RL", "RL", ...
## $ LotFrontage <int> 65, 80, 68, 60, 84, 85, 75, NA, 51, 50, 70, 85, ...
## $ LotArea <int> 8450, 9600, 11250, 9550, 14260, 14115, 10084, 10...
## $ Street <chr> "Pave", "Pave", "Pave", "Pave", "Pave", "Pave", "Pave", ...
## $ Alley <chr> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, ...
## $ LotShape <chr> "Reg", "Reg", "IR1", "IR1", "IR1", "IR1", "Reg",...
## $ LandContour <chr> "Lvl", "Lvl", "Lvl", "Lvl", "Lvl", "Lvl", "Lvl",...
## $ Utilities <chr> "AllPub", "AllPub", "AllPub", "AllPub", "AllPub"...
## $ LotConfig <chr> "Inside", "FR2", "Inside", "Corner", "FR2", "Ins...
## $ LandSlope <chr> "Gtl", "Gtl", "Gtl", "Gtl", "Gtl", "Gtl", "Gtl",...
## $ Neighborhood <chr> "CollgCr", "Veenker", "CollgCr", "Crawfor", "NoR...
## $ Condition1 <chr> "Norm", "Feedr", "Norm", "Norm", "Norm", "Norm",...
## $ Condition2 <chr> "Norm", "Norm", "Norm", "Norm", "Norm", "Norm", ...
## $ BldgType <chr> "1Fam", "1Fam", "1Fam", "1Fam", "1Fam", "1Fam", ...
## $ HouseStyle <chr> "2Story", "1Story", "2Story", "2Story", "2Story"...
## $ OverallQual <int> 7, 6, 7, 7, 8, 5, 8, 7, 7, 5, 5, 9, 5, 7, 6, 7, ...
## $ OverallCond <int> 5, 8, 5, 5, 5, 5, 5, 6, 5, 6, 5, 5, 6, 5, 5, 8, ...
## $ YearBuilt <int> 2003, 1976, 2001, 1915, 2000, 1993, 2004, 1973, ...
## $ YearRemodAdd <int> 2003, 1976, 2002, 1970, 2000, 1995, 2005, 1973, ...
## $ RoofStyle <chr> "Gable", "Gable", "Gable", "Gable", "Gable", "Ga...
## $ RoofMatl <chr> "CompShg", "CompShg", "CompShg", "CompShg", "Com...
## $ Exterior1st <chr> "VinylSd", "MetalSd", "VinylSd", "Wd Sdng", "Vin...
## $ Exterior2nd <chr> "VinylSd", "MetalSd", "VinylSd", "Wd Shng", "Vin...
## $ MasVnrType <chr> "BrkFace", "None", "BrkFace", "None", "BrkFace",...
## $ MasVnrArea <int> 196, 0, 162, 0, 350, 0, 186, 240, 0, 0, 0, 286, ...
## $ ExterQual <chr> "Gd", "TA", "Gd", "TA", "Gd", "TA", "Gd", "TA", ...
## $ ExterCond <chr> "TA", "TA", "TA", "TA", "TA", "TA", "TA", "TA", ...
## $ Foundation <chr> "PConc", "CBlock", "PConc", "BrkTil", "PConc", "...
## $ BsmtQual <chr> "Gd", "Gd", "Gd", "TA", "Gd", "Gd", "Ex", "Gd", ...
## $ BsmtCond <chr> "TA", "TA", "TA", "Gd", "TA", "TA", "TA", "TA", ...
## $ BsmtExposure <chr> "No", "Gd", "Mn", "No", "Av", "No", "Av", "Mn", ...
## $ BsmtFinType1 <chr> "GLQ", "ALQ", "GLQ", "ALQ", "GLQ", "GLQ", "GLQ",...
## $ BsmtFinSF1 <int> 706, 978, 486, 216, 655, 732, 1369, 859, 0, 851,...
## $ BsmtFinType2 <chr> "Unf", "Unf", "Unf", "Unf", "Unf", "Unf", "Unf",...
## $ BsmtFinSF2 <int> 0, 0, 0, 0, 0, 0, 0, 32, 0, 0, 0, 0, 0, 0, 0, ...
## $ BsmtUnfSF <int> 150, 284, 434, 540, 490, 64, 317, 216, 952, 140,...
## $ TotalBsmtSF <int> 856, 1262, 920, 756, 1145, 796, 1686, 1107, 952,...
## $ Heating <chr> "GasA", "GasA", "GasA", "GasA", "GasA", "GasA", ...
## $ HeatingQC <chr> "Ex", "Ex", "Ex", "Gd", "Ex", "Ex", "Ex", "Ex", ...
## $ CentralAir <chr> "Y", "Y", "Y", "Y", "Y", "Y", "Y", "Y", "Y", "Y"...
## $ Electrical <chr> "SBrkr", "SBrkr", "SBrkr", "SBrkr", "SBrkr", "SB...
## $ `1stFlrSF` <int> 856, 1262, 920, 961, 1145, 796, 1694, 1107, 1022...
## $ `2ndFlrSF` <int> 854, 0, 866, 756, 1053, 566, 0, 983, 752, 0, 0, ...
## $ LowQualFinSF <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
## $ GrLivArea <int> 1710, 1262, 1786, 1717, 2198, 1362, 1694, 2090, ...
## $ BsmtFullBath <int> 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, ...
## $ BsmtHalfBath <int> 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
```

```
## $ FullBath      <int> 2, 2, 2, 1, 2, 1, 2, 2, 2, 1, 1, 3, 1, 2, 1, 1, ...
## $ HalfBath      <int> 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, ...
## $ BedroomAbvGr <int> 3, 3, 3, 3, 4, 1, 3, 3, 2, 2, 3, 4, 2, 3, 2, 2, ...
## $ KitchenAbvGr <int> 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 1, 1, 1, 1, 1, ...
## $ KitchenQual   <chr> "Gd", "TA", "Gd", "Gd", "Gd", "TA", "Gd", "TA", ...
## $ TotRmsAbvGrd <int> 8, 6, 6, 7, 9, 5, 7, 7, 8, 5, 5, 11, 4, 7, 5, 5, ...
## $ Functional    <chr> "Typ", "Typ", "Typ", "Typ", "Typ", "Typ", "Typ", ...
## $ Fireplaces    <int> 0, 1, 1, 1, 1, 0, 1, 2, 2, 2, 0, 2, 0, 1, 1, 0, ...
## $ FireplaceQu   <chr> NA, "TA", "TA", "Gd", "TA", NA, "Gd", "TA", "TA"...
## $ GarageType    <chr> "Attchd", "Attchd", "Attchd", "Detchd", "Attchd"...
## $ GarageYrBlt   <int> 2003, 1976, 2001, 1998, 2000, 1993, 2004, 1973, ...
## $ GarageFinish  <chr> "RFn", "RFn", "RFn", "Unf", "RFn", "Unf", "RFn", ...
## $ GarageCars    <int> 2, 2, 2, 3, 3, 2, 2, 2, 2, 1, 1, 3, 1, 3, 1, 2, ...
## $ GarageArea    <int> 548, 460, 608, 642, 836, 480, 636, 484, 468, 205...
## $ GarageQual    <chr> "TA", "TA", "TA", "TA", "TA", "TA", "TA", "TA", ...
## $ GarageCond    <chr> "TA", "TA", "TA", "TA", "TA", "TA", "TA", "TA", ...
## $ PavedDrive    <chr> "Y", "Y", "Y", "Y", "Y", "Y", "Y", "Y", "Y", "Y", "Y"...
## $ WoodDeckSF    <int> 0, 298, 0, 0, 192, 40, 255, 235, 90, 0, 0, 147, ...
## $ OpenPorchSF   <int> 61, 0, 42, 35, 84, 30, 57, 204, 0, 4, 0, 21, 0, ...
## $ EnclosedPorch <int> 0, 0, 0, 272, 0, 0, 0, 228, 205, 0, 0, 0, 0, 0, ...
## $ `3SsnPorch`   <int> 0, 0, 0, 0, 0, 320, 0, 0, 0, 0, 0, 0, 0, 0, 0...
## $ ScreenPorch   <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 176, 0, 0, 0...
## $ PoolArea      <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
## $ PoolQC        <chr> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, ...
## $ Fence         <chr> NA, NA, NA, NA, NA, "MnPrv", NA, NA, NA, NA, NA, ...
## $ MiscFeature    <chr> NA, NA, NA, NA, NA, "Shed", NA, "Shed", NA, NA, ...
## $ MiscVal       <int> 0, 0, 0, 0, 0, 700, 0, 350, 0, 0, 0, 0, 0, 0, ...
## $ MoSold        <int> 2, 5, 9, 2, 12, 10, 8, 11, 4, 1, 2, 7, 9, 8, 5, ...
## $ YrSold        <int> 2008, 2007, 2008, 2006, 2008, 2009, 2007, 2009, ...
## $ SaleType       <chr> "WD", "WD", "WD", "WD", "WD", "WD", "WD", "WD", ...
## $ SaleCondition  <chr> "Normal", "Normal", "Normal", "Abnorml", "Normal..."
## $ SalePrice      <int> 208500, 181500, 223500, 140000, 250000, 143000, ...
```

So there are 1460 observations with 81 columns (variables)

Split the Data to training set and test set

Exercise 2

Normally at this point you would spend a few days on EDA, but for this homework we will get right to fitting some linear regression models.

Our first step is to randomly split the data into train and test datasets. We will use a 70/30 split. There is an R package that will do the split for you, but let's get some more practice with R and do it ourselves by filling in the blanks in the code below.

```
# load packages
library(tidyverse)
library(broom)

# When taking a random sample, it is often useful to set a seed so that
# your work is reproducible. Setting a seed will guarantee that the same
# random sample will be generated every time, so long as you always set the
# same seed beforehand
set.seed(29283)

# This data already has an Id column which we can make use of.
```

```
# Let's create our training set using sample_frac. Fill in the blank.
train_set <- train %>% sample_frac(____)
# let's create our testing set using the Id column. Fill in the blanks.
test_set <- train %>% filter(!(____ %in% ____$Id))
```

Answer

Let's fill in the blanks.

```
# When taking a random sample, it is often useful to set a seed so that
# your work is reproducible. Setting a seed will guarantee that the same
# random sample will be generated every time, so long as you always set the
# same seed beforehand
set.seed(29283)
# This data already has an Id column which we can make use of.
# Let's create our training set using sample_frac. Fill in the blank.
train_set <- train %>% sample_frac(0.7)
# Print train set
train_set
```

```
## # A tibble: 1,022 x 81
##       Id MSSubClass MSZoning LotFrontage LotArea Street Alley LotShape
##   <int>      <int> <chr>          <int>    <int> <chr>  <chr> <chr>
## 1    22         45 RM              57     7449 Pave   Grv1   Reg
## 2   637         30 RM              51     6120 Pave   <NA>   Reg
## 3   121         80 RL              NA    21453 Pave   <NA>   IR1
## 4   575         80 RL              70    10500 Pave   <NA>   Reg
## 5  1423        120 RM              37     4435 Pave   <NA>   Reg
## 6  1169         70 RL             120    13728 Pave   <NA>   Reg
## 7  1261         60 RL              NA    24682 Pave   <NA>   IR3
## 8  1319         20 RL              NA    14781 Pave   <NA>   IR2
## 9   116        160 FV              34     3230 Pave   Pave   Reg
## 10 1125         80 RL              NA     9125 Pave   <NA>   IR1
## # ... with 1,012 more rows, and 73 more variables: LandContour <chr>,
## #   Utilities <chr>, LotConfig <chr>, LandSlope <chr>, Neighborhood <chr>,
## #   Condition1 <chr>, Condition2 <chr>, BldgType <chr>, HouseStyle <chr>,
## #   OverallQual <int>, OverallCond <int>, YearBuilt <int>,
## #   YearRemodAdd <int>, RoofStyle <chr>, RoofMatl <chr>,
## #   Exterior1st <chr>, Exterior2nd <chr>, MasVnrType <chr>,
## #   MasVnrArea <int>, ExterQual <chr>, ExterCond <chr>, Foundation <chr>,
## #   BsmtQual <chr>, BsmtCond <chr>, BsmtExposure <chr>,
## #   BsmtFinType1 <chr>, BsmtFinSF1 <int>, BsmtFinType2 <chr>,
## #   BsmtFinSF2 <int>, BsmtUnfSF <int>, TotalBsmtSF <int>, Heating <chr>,
## #   HeatingQC <chr>, CentralAir <chr>, Electrical <chr>, `1stFlrSF` <int>,
## #   `2ndFlrSF` <int>, LowQualFinSF <int>, GrLivArea <int>,
## #   BsmtFullBath <int>, BsmtHalfBath <int>, FullBath <int>,
## #   HalfBath <int>, BedroomAbvGr <int>, KitchenAbvGr <int>,
## #   KitchenQual <chr>, TotRmsAbvGrd <int>, Functional <chr>,
## #   Fireplaces <int>, FireplaceQu <chr>, GarageType <chr>,
## #   GarageYrBlt <int>, GarageFinish <chr>, GarageCars <int>,
## #   GarageArea <int>, GarageQual <chr>, GarageCond <chr>,
## #   PavedDrive <chr>, WoodDeckSF <int>, OpenPorchSF <int>,
## #   EnclosedPorch <int>, `3SsnPorch` <int>, ScreenPorch <int>,
## #   PoolArea <int>, PoolQC <chr>, Fence <chr>, MiscFeature <chr>,
## #   MiscVal <int>, MoSold <int>, YrSold <int>, SaleType <chr>,
```

```
## # SaleCondition <chr>, SalePrice <int>
# let's create our testing set using the Id column. Fill in the blanks.
test_set <- train %>% filter(!(train$Id %in% train_set$Id))
# Print test set
test_set

## # A tibble: 438 x 81
##       Id MSSubClass MSZoning LotFrontage LotArea Street Alley LotShape
##   <int>      <int> <chr>          <int>    <int> <chr>  <chr> <chr>
## 1     1         60 RL             65     8450 Pave  <NA>  Reg
## 2     2         20 RL             80     9600 Pave  <NA>  Reg
## 3     3         60 RL             68    11250 Pave  <NA>  IR1
## 4     4         70 RL             60     9550 Pave  <NA>  IR1
## 5     14        20 RL             91    10652 Pave  <NA>  IR1
## 6     23        20 RL             75     9742 Pave  <NA>  Reg
## 7     27        20 RL             60     7200 Pave  <NA>  Reg
## 8     38        20 RL             74     8532 Pave  <NA>  Reg
## 9     40        90 RL             65     6040 Pave  <NA>  Reg
## 10    42        20 RL            115    16905 Pave  <NA>  Reg
## # ... with 428 more rows, and 73 more variables: LandContour <chr>,
## # Utilities <chr>, LotConfig <chr>, LandSlope <chr>, Neighborhood <chr>,
## # Condition1 <chr>, Condition2 <chr>, BldgType <chr>, HouseStyle <chr>,
## # OverallQual <int>, OverallCond <int>, YearBuilt <int>,
## # YearRemodAdd <int>, RoofStyle <chr>, RoofMatl <chr>,
## # Exterior1st <chr>, Exterior2nd <chr>, MasVnrType <chr>,
## # MasVnrArea <int>, ExterQual <chr>, ExterCond <chr>, Foundation <chr>,
## # BsmtQual <chr>, BsmtCond <chr>, BsmtExposure <chr>,
## # BsmtFinType1 <chr>, BsmtFinSF1 <int>, BsmtFinType2 <chr>,
## # BsmtFinSF2 <int>, BsmtUnfSF <int>, TotalBsmtSF <int>, Heating <chr>,
## # HeatingQC <chr>, CentralAir <chr>, Electrical <chr>, `1stFlrSF` <int>,
## # `2ndFlrSF` <int>, LowQualFinSF <int>, GrLivArea <int>,
## # BsmtFullBath <int>, BsmtHalfBath <int>, FullBath <int>,
## # HalfBath <int>, BedroomAbvGr <int>, KitchenAbvGr <int>,
## # KitchenQual <chr>, TotRmsAbvGrd <int>, Functional <chr>,
## # Fireplaces <int>, FireplaceQu <chr>, GarageType <chr>,
## # GarageYrBlt <int>, GarageFinish <chr>, GarageCars <int>,
## # GarageArea <int>, GarageQual <chr>, GarageCond <chr>,
## # PavedDrive <chr>, WoodDeckSF <int>, OpenPorchSF <int>,
## # EnclosedPorch <int>, `3SsnPorch` <int>, ScreenPorch <int>,
## # PoolArea <int>, PoolQC <chr>, Fence <chr>, MiscFeature <chr>,
## # MiscVal <int>, MoSold <int>, YrSold <int>, SaleType <chr>,
## # SaleCondition <chr>, SalePrice <int>
```

Now, we have separated our train data set and test data set.

The intercept: *SalePrice*

Exercise 3

Our target is called *SalePrice*. First, we can fit a simple regression model consisting of only the intercept (the average of *SalePrice*). Fit the model and then use the broom package to

- take a look at the coefficient,
- compare the coefficient to the average value of *SalePrice*, and

- take a look at the R-squared.

Use the code below and fill in the blanks.

```
# Fit a model with intercept only
mod_0 <- lm(SalePrice ~ 1, data = _____)
# Double-check that the average SalePrice is equal to our model's coefficient
mean(train_set$SalePrice)
tidy(____)
# Check the R-squared
glance(_____)
```

Answer

Let's fill in the blanks.

```
# Fit a model with intercept only
mod_0 <- lm(SalePrice ~ 1, data = train_set)
summary(mod_0)

##
## Call:
## lm(formula = SalePrice ~ 1, data = train_set)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -142876  -52251  -18181   32824  562824
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   182176      2492    73.1  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 79670 on 1021 degrees of freedom
# Double-check that the average SalePrice is equal to our model's coefficient
mean(train_set$SalePrice)

## [1] 182176
tidy(mod_0)

##           term estimate std.error statistic p.value
## 1 (Intercept)   182176   2492.072    73.10222      0
# Check the R-squared
glance(mod_0)

##   r.squared adj.r.squared   sigma statistic p.value df   logLik   AIC
## 1         0             0 79668.37      NA      NA  1 -12983.57 25971.13
##           BIC      deviance df.residual
## 1 25980.99 6.480338e+12          1021
```

EDA on *GrLivArea*, *OverallQual*, and *Neighborhood*

Exercise 4

Now fit a linear regression model using GrLivArea, OverallQual, and Neighborhood as the features. Don't forget to look at data_description.txt to understand what these variables mean. Ask yourself these questions before fitting the model:

- What kind of relationship will these features have with our target?
- Can the relationship be estimated linearly?
- Are these good features, given the problem we are trying to solve?

After fitting the model, output the coefficients and the R-squared using the broom package.

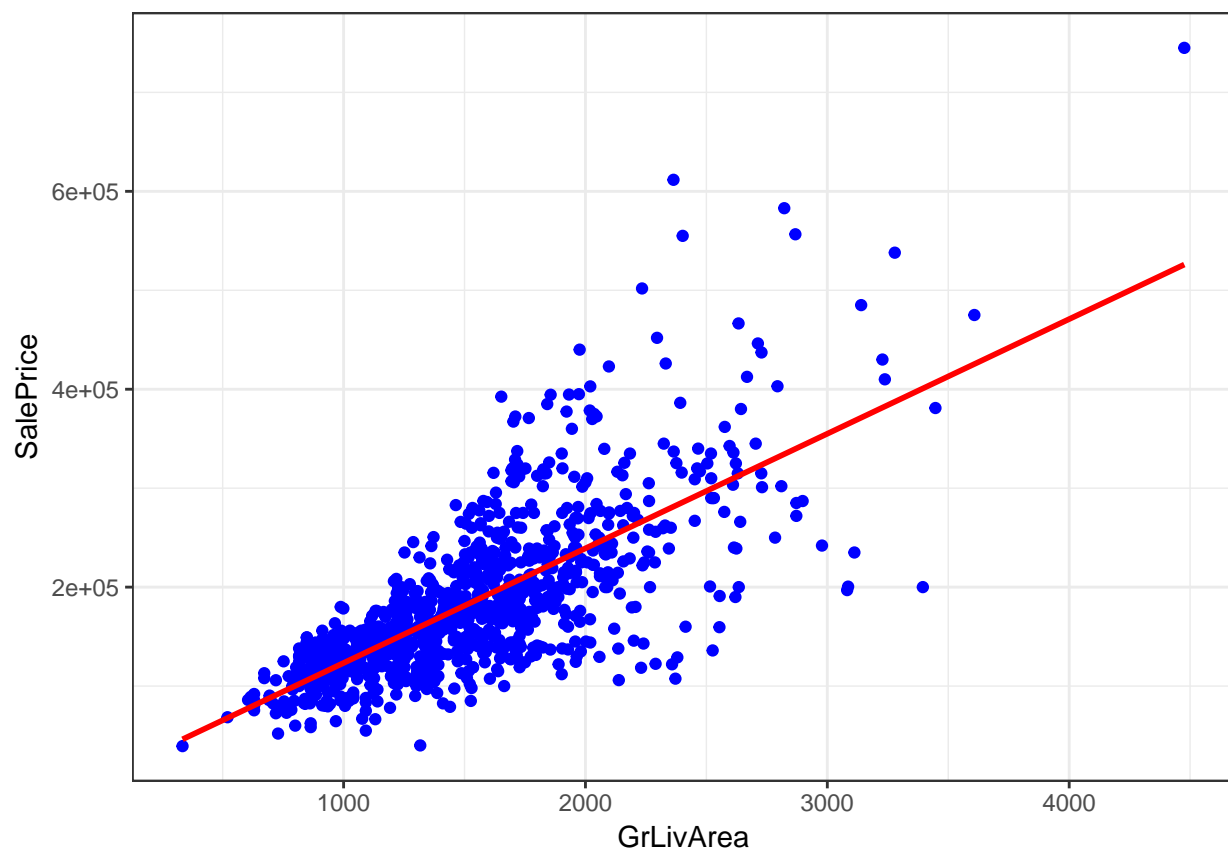
Answer these questions:

- How would you interpret the coefficients on GrLivArea and OverallQual?
- How would you interpret the coefficient on NeighborhoodBrkSide?
- Are the features significant?
- Are the features practically significant?
- Is the model a good fit (to the training set)?

Answer

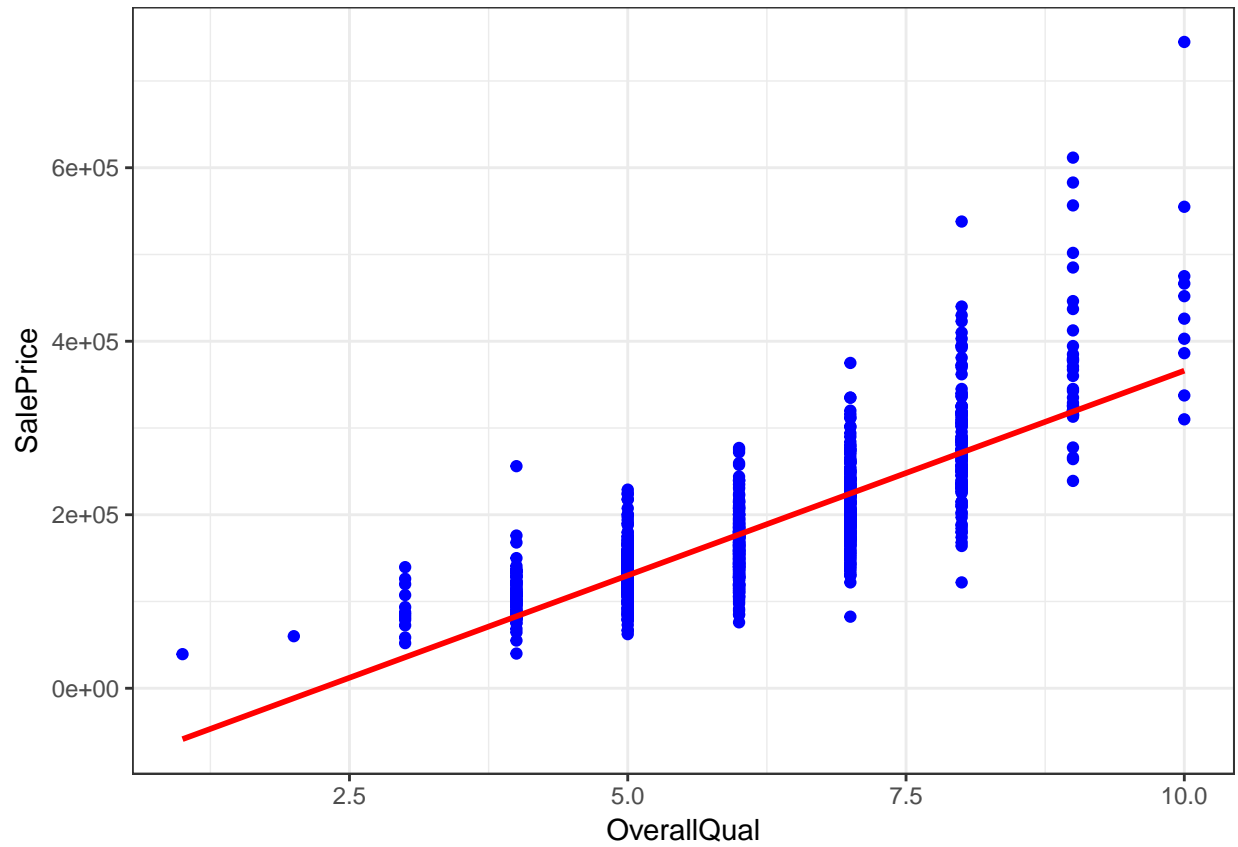
Let's plot the graph to see the relationship between SalePrice and GrLivArea, OverallQual, and Neighborhood.

```
ggplot(train_set, aes(x = GrLivArea, y = SalePrice)) +  
  geom_point(color = "blue") +  
  geom_smooth(method = "lm", se = FALSE, color = "red") +  
  theme_bw()
```

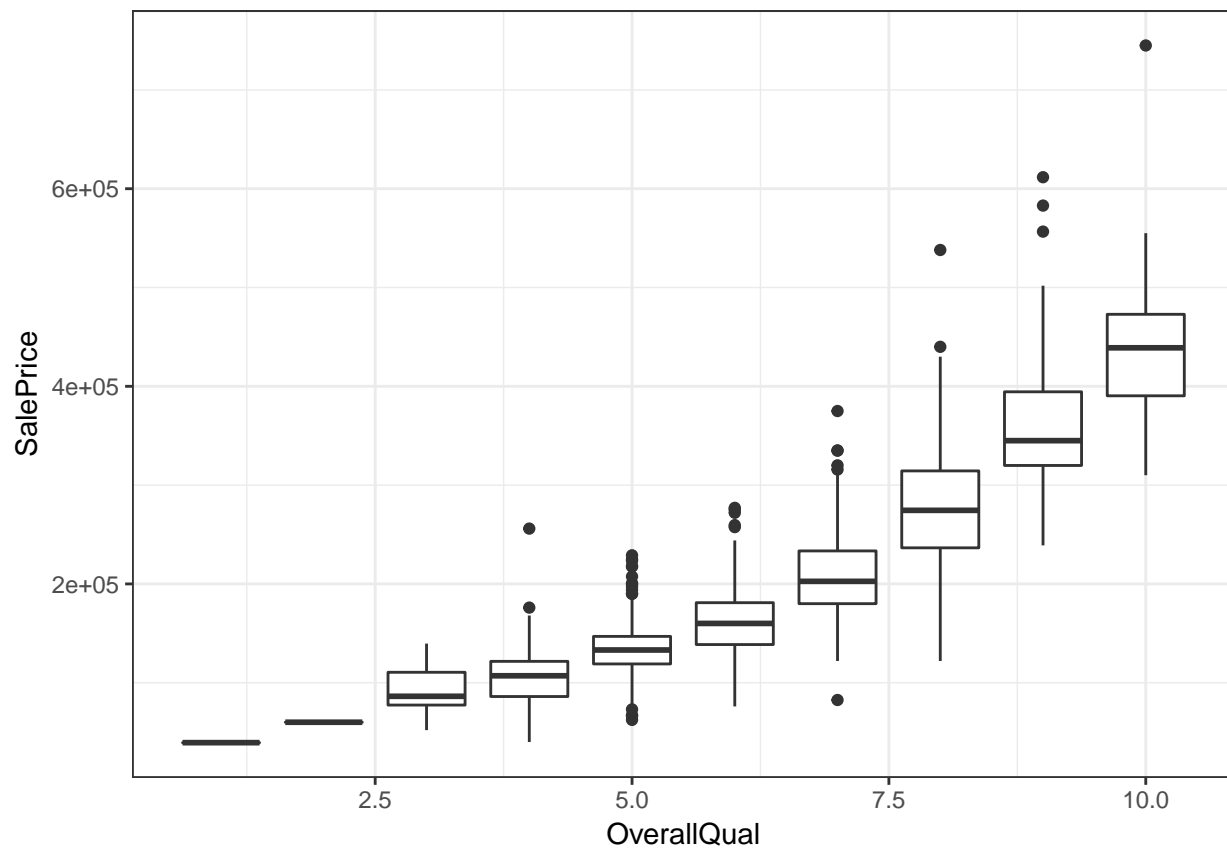


The relationship between SalePrice and GrLivArea is linear. It can be estimated linearly. It is an important feature to estimate as it impacts the SalePrice.


```
ggplot(train_set, aes(x = OverallQual, y = SalePrice)) +
  geom_point(color = "blue") +
  geom_smooth(method = "lm", se = FALSE, color = "red") +
  theme_bw()
```




```
ggplot(train_set, aes(x = OverallQual, y = SalePrice, group=OverallQual)) +
  geom_boxplot() +
  theme_bw()
```



The relationship between SalePrice and OverallQual is linear. It can be estimated linearly. It is an important feature to estimate as it impacts the SalePrice.

```
train_set %>%
ggplot() +
  geom_bar(aes(x = Neighborhood, y = SalePrice), stat = 'identity') +
  theme_bw()
```



There is a relationship between SalePrice and Neighborhood. But the Neighborhood is categorical, so no linear relationship with SalePrice. It is an important feature to estimate as it impacts the SalePrice. We will factor it and add to our regression model.

Let's create the models and see the values.

```
# Fit a model for GrLivArea
lm_1 <- lm(SalePrice ~ GrLivArea, data = train_set)
mean(train_set$SalePrice)

## [1] 182176

summary(lm_1)

##
## Call:
## lm(formula = SalePrice ~ GrLivArea, data = train_set)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -200772  -29953    -654    22765   330309
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  7518.567   5334.587   1.409   0.159
## GrLivArea     115.833     3.355   34.526 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## Residual standard error: 54130 on 1020 degrees of freedom
## Multiple R-squared:  0.5389, Adjusted R-squared:  0.5384
## F-statistic: 1192 on 1 and 1020 DF,  p-value: < 2.2e-16

tidy(lm_1)

##           term estimate  std.error statistic      p.value
## 1 (Intercept) 7518.567 5334.586903   1.40940 1.590217e-01
## 2   GrLivArea  115.833   3.354991  34.52558 1.189915e-173

glance(lm_1)

##   r.squared adj.r.squared   sigma statistic      p.value df logLik   AIC
## 1 0.5388821    0.5384301 54125.85  1192.016 1.189915e-173  2 -12588 25182
##           BIC   deviance df.residual
## 1 25196.79 2.9882e+12      1020

# Fit a model for OverallQual
lm_2 <- lm(SalePrice ~ OverallQual, data = train_set)
mean(train_set$SalePrice)

## [1] 182176

summary(lm_2)

##
## Call:
## lm(formula = SalePrice ~ OverallQual, data = train_set)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -149666  -29233   -1589    20175   378950
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -105873     6894   -15.36  <2e-16 ***
## OverallQual    47192     1103    42.80  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 47670 on 1020 degrees of freedom
## Multiple R-squared:  0.6423, Adjusted R-squared:  0.642
## F-statistic: 1832 on 1 and 1020 DF,  p-value: < 2.2e-16

tidy(lm_2)

##           term estimate std.error statistic      p.value
## 1 (Intercept) -105872.5  6893.466  -15.35839 4.801210e-48
## 2 OverallQual   47192.3  1102.649   42.79902 5.918164e-230

glance(lm_2)

##   r.squared adj.r.squared   sigma statistic      p.value df   logLik
## 1 0.6423257    0.641975 47669.72  1831.756 5.918164e-230  2 -12458.19
##           AIC      BIC   deviance df.residual
## 1 24922.38 24937.17 2.317851e+12      1020
```

```
# Fit a model for Neighborhood
train_set <- train_set %>% mutate(Neighborhood_fct = factor(Neighborhood, ordered = FALSE))
lm_3 <- lm(SalePrice ~ Neighborhood_fct, data = train_set)
mean(train_set$SalePrice)
```

```
## [1] 182176
```

```
summary(lm_3)
```

```
##
## Call:
## lm(formula = SalePrice ~ Neighborhood_fct, data = train_set)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -164344  -27697   -5142   19353  409437
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      200309      14880  13.461 < 2e-16 ***
## Neighborhood_fctBlueste -76309      55677  -1.371  0.170820
## Neighborhood_fctBrDale -92909      21980  -4.227  2.59e-05 ***
## Neighborhood_fctBrkSide -76713      16813  -4.563  5.68e-06 ***
## Neighborhood_fctClearCr  19854      20330   0.977  0.329026
## Neighborhood_fctCollgCr  -4492      15774  -0.285  0.775900
## Neighborhood_fctCrawfor  4905      17077   0.287  0.773989
## Neighborhood_fctEdwards -69405      16300  -4.258  2.26e-05 ***
## Neighborhood_fctGilbert  -9248      16490  -0.561  0.575045
## Neighborhood_fctIDOTRR  -94072      18769  -5.012  6.37e-07 ***
## Neighborhood_fctMeadowV -103085      21044  -4.899  1.13e-06 ***
## Neighborhood_fctMitchel -44659      17728  -2.519  0.011920 *
## Neighborhood_fctNAMES   -54161      15455  -3.504  0.000478 ***
## Neighborhood_fctNoRidge  135254      18617   7.265  7.50e-13 ***
## Neighborhood_fctNPkVill -59059      25152  -2.348  0.019067 *
## Neighborhood_fctNridgHt  118035      16438   7.181  1.36e-12 ***
## Neighborhood_fctNWAMES  -12272      16637  -0.738  0.460904
## Neighborhood_fctOldTown -72381      16134  -4.486  8.10e-06 ***
## Neighborhood_fctSawyer  -63193      16703  -3.783  0.000164 ***
## Neighborhood_fctSawyerW -20426      16981  -1.203  0.229323
## Neighborhood_fctSomerst  34566      16413   2.106  0.035457 *
## Neighborhood_fctStoneBr  118544      19311   6.139  1.20e-09 ***
## Neighborhood_fctSWISU   -52975      20033  -2.644  0.008313 **
## Neighborhood_fctTimber   52346      18225   2.872  0.004162 **
## Neighborhood_fctVeenker  63620      25152   2.529  0.011579 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 53650 on 997 degrees of freedom
## Multiple R-squared:  0.5571, Adjusted R-squared:  0.5465
## F-statistic: 52.26 on 24 and 997 DF,  p-value: < 2.2e-16
```

```
tidy(lm_3)
```

```
##              term      estimate std.error  statistic      p.value
## 1      (Intercept) 200308.462  14880.25  13.4613631 4.487217e-38
```

```
## 2 Neighborhood_fctBlueste -76308.462 55676.80 -1.3705612 1.708203e-01
## 3 Neighborhood_fctBrDale -92908.462 21979.59 -4.2270339 2.585277e-05
## 4 Neighborhood_fctBrkSide -76713.249 16812.68 -4.5628209 5.675519e-06
## 5 Neighborhood_fctClearCr 19853.672 20330.29 0.9765561 3.290259e-01
## 6 Neighborhood_fctCollgCr -4491.700 15774.54 -0.2847437 7.758997e-01
## 7 Neighborhood_fctCrawfor 4905.221 17077.14 0.2872390 7.739890e-01
## 8 Neighborhood_fctEdwards -69405.077 16300.50 -4.2578500 2.259005e-05
## 9 Neighborhood_fctGilbert -9247.988 16490.05 -0.5608225 5.750446e-01
## 10 Neighborhood_fctIDOTRR -94072.098 18768.65 -5.0121942 6.365074e-07
## 11 Neighborhood_fctMeadowV -103085.385 21043.85 -4.8985984 1.125722e-06
## 12 Neighborhood_fctMitchel -44658.462 17727.84 -2.5191151 1.192047e-02
## 13 Neighborhood_fctNAMES -54161.231 15455.33 -3.5043723 4.780249e-04
## 14 Neighborhood_fctNoRidge 135254.147 18616.48 7.2652908 7.499696e-13
## 15 Neighborhood_fctNPkVill -59058.462 25152.21 -2.3480422 1.906702e-02
## 16 Neighborhood_fctNridgHt 118035.386 16438.06 7.1806164 1.355530e-12
## 17 Neighborhood_fctNWAmes -12271.923 16636.63 -0.7376449 4.609038e-01
## 18 Neighborhood_fctOldTown -72380.651 16134.44 -4.4860949 8.097746e-06
## 19 Neighborhood_fctSawyer -63192.622 16703.04 -3.7833005 1.639801e-04
## 20 Neighborhood_fctSawyerW -20425.787 16981.27 -1.2028425 2.293228e-01
## 21 Neighborhood_fctSomerst 34565.772 16413.31 2.1059604 3.545731e-02
## 22 Neighborhood_fctStoneBr 118543.907 19311.16 6.1386206 1.199918e-09
## 23 Neighborhood_fctSWISU -52975.087 20033.15 -2.6443711 8.312879e-03
## 24 Neighborhood_fctTimber 52345.885 18224.51 2.8722792 4.161598e-03
## 25 Neighborhood_fctVeenker 63620.110 25152.21 2.5294039 1.157881e-02
```

```
glance(lm_3)
```

```
## r.squared adj.r.squared sigma statistic p.value df logLik
## 1 0.5571452 0.5464847 53651.51 52.26259 1.011167e-157 25 -12567.35
## AIC BIC deviance df.residual
## 1 25186.7 25314.87 2.869849e+12 997
```

```
# Fit a model with all three variables
```

```
lm_4 <- lm(SalePrice ~ GrLivArea + OverallQual + Neighborhood_fct, data = train_set)
mean(train_set$SalePrice)
```

```
## [1] 182176
```

```
summary(lm_4)
```

```
##
## Call:
## lm(formula = SalePrice ~ GrLivArea + OverallQual + Neighborhood_fct,
## data = train_set)
##
## Residuals:
## Min 1Q Median 3Q Max
## -118777 -17495 -57 16149 249166
##
## Coefficients:
## Estimate Std. Error t value Pr(>|t|)
## (Intercept) -45017.875 12933.342 -3.481 0.000522 ***
## GrLivArea 62.777 3.006 20.884 < 2e-16 ***
## OverallQual 21692.232 1353.714 16.024 < 2e-16 ***
## Neighborhood_fctBlueste -38288.881 36531.907 -1.048 0.294850
## Neighborhood_fctBrDale -43314.054 14524.694 -2.982 0.002933 **
```

```
## Neighborhood_fctBrkSide -14064.371 11318.850 -1.243 0.214322
## Neighborhood_fctClearCr 27839.007 13561.347 2.053 0.040351 *
## Neighborhood_fctCollgCr 4297.674 10372.304 0.414 0.678713
## Neighborhood_fctCrawfor 7423.056 11371.512 0.653 0.514051
## Neighborhood_fctEdwards -15284.115 10994.287 -1.390 0.164783
## Neighborhood_fctGilbert -8357.559 10894.173 -0.767 0.443169
## Neighborhood_fctIDOTRR -32689.431 12603.713 -2.594 0.009636 **
## Neighborhood_fctMeadowV -14446.065 14190.149 -1.018 0.308909
## Neighborhood_fctMitchel 1922.315 11788.608 0.163 0.870500
## Neighborhood_fctNAMES -7719.679 10375.956 -0.744 0.457054
## Neighborhood_fctNoRidge 47685.168 12567.433 3.794 0.000157 ***
## Neighborhood_fctNPkVill -20240.711 16548.665 -1.223 0.221581
## Neighborhood_fctNridgHt 63872.808 10880.457 5.870 5.92e-09 ***
## Neighborhood_fctNWames -12279.333 11047.503 -1.112 0.266620
## Neighborhood_fctOldTown -36107.076 10849.171 -3.328 0.000906 ***
## Neighborhood_fctSawyer -4121.925 11252.370 -0.366 0.714207
## Neighborhood_fctSawyerW -5391.971 11230.758 -0.480 0.631257
## Neighborhood_fctSomerst 18700.967 10772.213 1.736 0.082867 .
## Neighborhood_fctStoneBr 65712.459 12745.313 5.156 3.05e-07 ***
## Neighborhood_fctSWISU -45451.867 13564.793 -3.351 0.000836 ***
## Neighborhood_fctTimber 27925.086 11985.326 2.330 0.020009 *
## Neighborhood_fctVeenker 54913.128 16521.075 3.324 0.000920 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 35180 on 995 degrees of freedom
## Multiple R-squared: 0.81, Adjusted R-squared: 0.805
## F-statistic: 163.1 on 26 and 995 DF, p-value: < 2.2e-16
```

```
tidy(lm_4)
```

```
##           term      estimate  std.error statistic
## 1 (Intercept) -45017.87483 12933.341808 -3.4807612
## 2 GrLivArea    62.77735    3.006033 20.8837885
## 3 OverallQual 21692.23178 1353.714104 16.0242342
## 4 Neighborhood_fctBlueste -38288.88063 36531.907177 -1.0480942
## 5 Neighborhood_fctBrDale -43314.05372 14524.693991 -2.9820975
## 6 Neighborhood_fctBrkSide -14064.37052 11318.850018 -1.2425618
## 7 Neighborhood_fctClearCr 27839.00662 13561.346871 2.0528202
## 8 Neighborhood_fctCollgCr 4297.67432 10372.304467 0.4143413
## 9 Neighborhood_fctCrawfor 7423.05573 11371.511784 0.6527765
## 10 Neighborhood_fctEdwards -15284.11495 10994.287187 -1.3901870
## 11 Neighborhood_fctGilbert -8357.55930 10894.173472 -0.7671586
## 12 Neighborhood_fctIDOTRR -32689.43085 12603.712743 -2.5936350
## 13 Neighborhood_fctMeadowV -14446.06504 14190.148622 -1.0180348
## 14 Neighborhood_fctMitchel 1922.31487 11788.608170 0.1630655
## 15 Neighborhood_fctNAMES -7719.67883 10375.956174 -0.7439969
## 16 Neighborhood_fctNoRidge 47685.16790 12567.432633 3.7943444
## 17 Neighborhood_fctNPkVill -20240.71145 16548.664867 -1.2231024
## 18 Neighborhood_fctNridgHt 63872.80848 10880.456671 5.8704161
## 19 Neighborhood_fctNWames -12279.33299 11047.502893 -1.1115030
## 20 Neighborhood_fctOldTown -36107.07577 10849.170903 -3.3280954
## 21 Neighborhood_fctSawyer -4121.92502 11252.369778 -0.3663162
## 22 Neighborhood_fctSawyerW -5391.97074 11230.758221 -0.4801075
## 23 Neighborhood_fctSomerst 18700.96725 10772.212794 1.7360377
```

```
## 24 Neighborhood_fctStoneBr 65712.45881 12745.312907 5.1558137
## 25 Neighborhood_fctSWISU -45451.86707 13564.792586 -3.3507233
## 26 Neighborhood_fctTimber 27925.08619 11985.325857 2.3299397
## 27 Neighborhood_fctVeenker 54913.12768 16521.075497 3.3238228
##      p.value
## 1 5.216927e-04
## 2 1.337222e-80
## 3 1.389020e-51
## 4 2.948497e-01
## 5 2.932566e-03
## 6 2.143221e-01
## 7 4.035110e-02
## 8 6.787135e-01
## 9 5.140512e-01
## 10 1.647830e-01
## 11 4.431692e-01
## 12 9.636216e-03
## 13 3.089089e-01
## 14 8.705000e-01
## 15 4.570540e-01
## 16 1.569690e-04
## 17 2.215806e-01
## 18 5.917964e-09
## 19 2.666204e-01
## 20 9.064637e-04
## 21 7.142070e-01
## 22 6.312565e-01
## 23 8.286672e-02
## 24 3.045915e-07
## 25 8.363074e-04
## 26 2.000859e-02
## 27 9.203087e-04
```

```
glance(lm_4)
```

```
##      r.squared adj.r.squared   sigma statistic p.value df    logLik      AIC
## 1 0.8099927      0.8050277 35178.1  163.1401      0 27 -12134.95 24325.91
##      BIC      deviance df.residual
## 1 24463.93 1.231311e+12          995
```

How would you interpret the coefficients on GrLivArea and OverallQual?

The coefficient for the GrLivArea predictor is 115.833. This means that for every increase by one square foot the house price increases by 115.833 dollars.

The coefficient for the OverallQual predictor is 47192. This means that for every increase by one point for the overall quality, the house price increases by 47192 dollars.

Answer

How would you interpret the coefficient on NeighborhoodBrkSide?

Answer

Are the features significant?

Answer

Are the features practically significant?

Answer

Is the model a good fit (to the training set)?

Answer

The adjusted R2 for the model is 0.8050277. This means that the model explains 80.5% of variability of the response data around its mean. I would say it is not a best model. We have to add more variables to get higher accuracy.

Evaluate the model

Exercise 5

Evaluate the model on test_set using the root mean squared error (RMSE). Use the predict function to get the model predictions for the testing set.

Hint: use the sqrt() and mean() functions:

```
test_predictions <- predict(NAME_OF_YOUR_MODEL_HERE, newdata = test_set)
rmse <- sqrt(mean((___ - ___)^2))
```

Answer

Let's predict on the test data and evaluate the model.

```
test_set <- test_set %>% mutate(Neighborhood_fct = factor(Neighborhood, ordered = FALSE))
test_predictions <- predict(lm_4, newdata = test_set)
rmse <- sqrt(mean((test_set$SalePrice - test_predictions)^2))
rmse
```

```
## [1] 41915.27
```

From the above RMSE value, we can conclude, maybe it is not a best model to predict the sales price.

Linear Model : downside

Exercise 7

One downside of the linear model is that it is sensitive to unusual values because the distance incorporates a squared term. Fit a linear model to the simulated data below, and visualise the results. Rerun a few times to generate different simulated datasets. What do you notice about the model?

```
sim1a <- tibble(
  x = rep(1:10, each = 3),
  y = x * 1.5 + 6 + rt(length(x), df = 2)
)
```


Answer

Let's create a model and run on the simulated data and visualize it.

```
sim1a <- tibble(
  x = rep(1:10, each = 3),
  y = x * 1.5 + 6 + rt(length(x), df = 2)
)

mod_5 <- lm(y~x, data = sim1a)
```

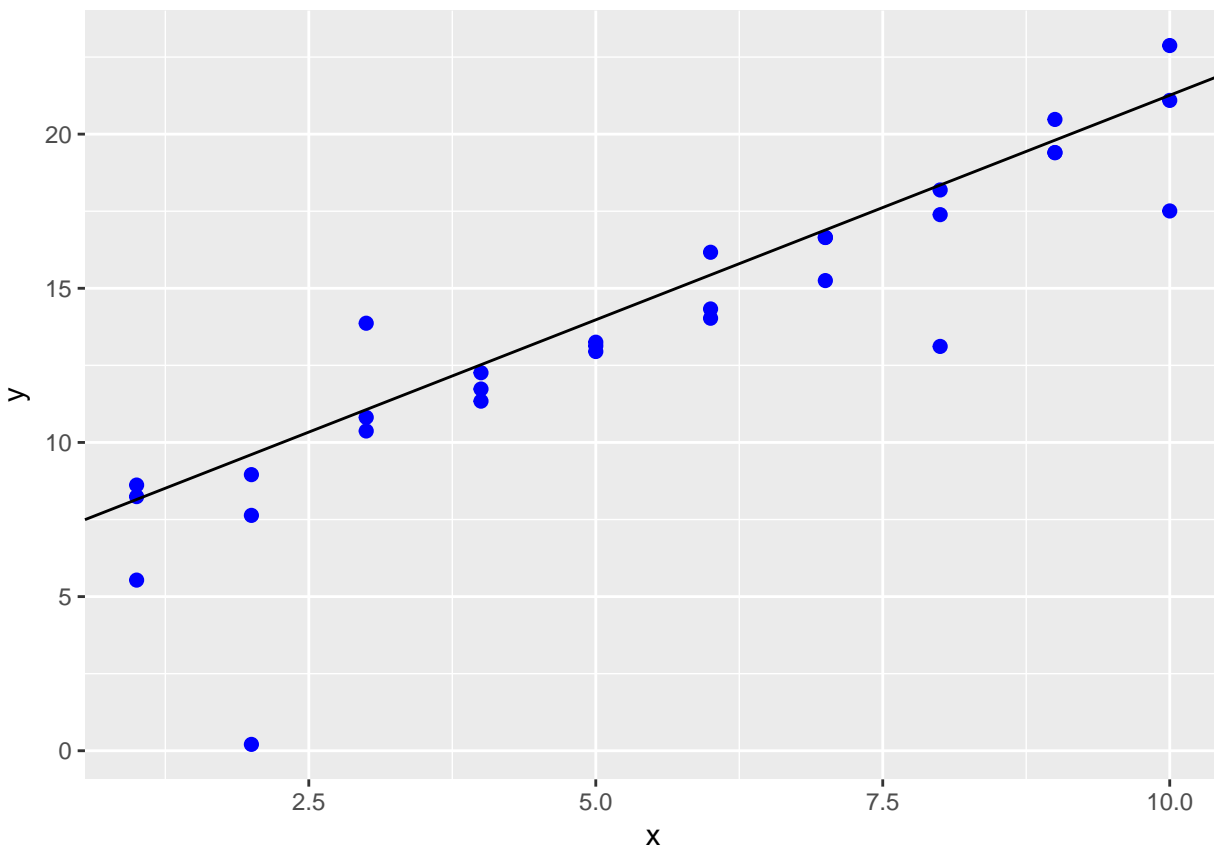
```
ggplot(sim1a,aes(x,y))+
  geom_point(size = 2, color = "blue")+
  geom_abline(intercept = mod_5$coefficients[1],slope = mod_5$coefficients[2])
```



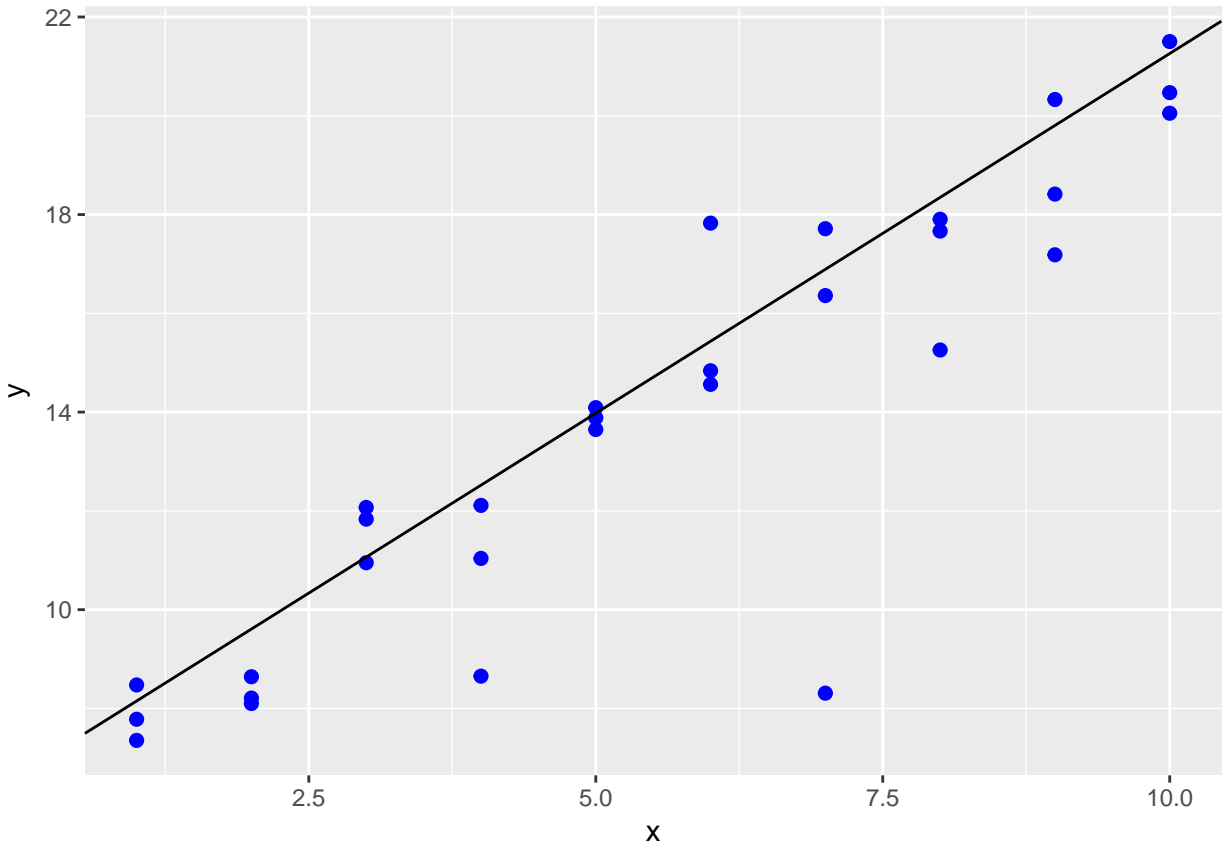
Now, let's run for few times.

```
sim2a <- tibble(
  x = rep(1:10, each = 3),
  y = x * 1.5 + 6 + rt(length(x), df = 2)
)

mod_5 <- lm(y~x, data = sim1a)
ggplot(sim2a,aes(x,y))+
  geom_point(size = 2, color = "blue")+
  geom_abline(intercept = mod_5$coefficients[1],slope = mod_5$coefficients[2])
```



```
sim3a <- tibble(  
  x = rep(1:10, each = 3),  
  y = x * 1.5 + 6 + rt(length(x), df = 2)  
)  
  
mod_5 <- lm(y~x, data = sim1a)  
ggplot(sim3a,aes(x,y))+  
  geom_point(size = 2, color = "blue")+  
  geom_abline(intercept = mod_5$coefficients[1],slope = mod_5$coefficients[2])
```



Conclusion: Sometimes, one single abnormal value forces the fitted line deviate from the “intuitively” best lines.

End of Homework 7