



COMPSCIX 415.2 Homework 5/Midterm

Sanatan Das

March 1, 2018

Contents

Code and Documents Git Repository	2
Load packages (prerequisites to run the code in this document)	2
RStudio and R Markdown (3 points)	2
The tidyverse packages (3 points)	2
R Basics (1.5 points)	3
Data import/export (3 points)	4
Visualization (6 points)	5
Data munging and wrangling (6 points)	7
EDA (6 points)	9
Git and Github (1.5 points)	16

Code and Documents Git Repository

All the work can be found in the below Git repository location:

<https://github.com/sanatanonline/compscix-415-2-assignments>

Load packages (prerequisites to run the code in this document)

```
library(tidyverse)
library(nycflights13)
library(dplyr)
```

RStudio and R Markdown (3 points)

1. Use markdown headers in your document to clearly separate each midterm question and add a table of contents to your document.

Answer

The following is the code used for markdown header and to add a table of contents in the document.

```
title: "COMPSCIX 415.2 Homework 5/Midterm"
author: "Sanatan Das"
date: "March 2, 2018"
output:
  html_document:
    number_sections: yes
    toc: yes
    toc_depth: 2
  pdf_document:
    toc: yes
    toc_depth: '2'
```

The tidyverse packages (3 points)

1. Can you name which package is associated with each task below?

- Plotting -
- Data munging/wrangling -
- Reshaping (speading and gathering) data -
- Importing/exporting data -

Answer:

The below are the packages associated with each task below.

- Plotting - Plotting is done mainly using *ggplot2* package which is a core member of *tidyverse* package.
- Data munging/wrangling - This mainly uses base R packages, *tibble* (which is a core member of *tidyverse* package) and *dplyr* package. We have used the datasets from *nycflights13* package.

- Reshaping (spreading and gathering) data - Reshaping of data is done using the functions of base R, *tibble* (which is a core member of *tidyverse* package) and *tidyr* packages.
- Importing/exporting data - Import and export of data mainly uses the functions of *readR* package which is a core member of *tidyverse* package.

2. Now can you name two functions that you’ve used from each package that you listed above for these tasks?

- Plotting -
- Data munging/wrangling -
- Reshaping data -
- Importing/exporting data (note that readRDS and saveRDS are base R functions) -

Answer:

The following are functions used from the packages listed above for the tasks.

- Plotting - *geom_point()* and *geom_smooth()*
- Data munging/wrangling - *filter()* and *arrange()*
- Reshaping data - *spread()* and *gather()*
- Importing/exporting data (note that readRDS and saveRDS are base R functions) - *read_csv()* and *write_csv()*

R Basics (1.5 points)

1. Fix this code with the fewest number of changes possible so it works:

```
My_data.name___is.too00ooLong! <- c( 1 , 2 , 3 )
```

Answer

```
My_data.name___is.too00ooLong <- c( 1 , 2 , 3 )
My_data.name___is.too00ooLong
```

```
## [1] 1 2 3
```

Explanation: ‘!’ is not allowed in a variable name. If the code is executed, R gives the below error:

Error: unexpected ‘!’ in “My_data.name___is.too00ooLong!”

2. Fix this code so it works:

```
my_string <- C('has', 'an', 'error', 'in', 'it')
```

Answer

There are two issues in the above code.

- `my_string <- C()`, in this code “C” is in uppercase whereas it should be lowercase. R is case sensitive.
- The last element is not enclosed by quote. So R can not parse it and throws parse error.

The correct code should be as below:

```
my_string <- c('has', 'an', 'error', 'in', 'it')
my_string
```

```
## [1] "has" "an" "error" "in" "it"
```

3. Look at the code below and comment on what happened to the values in the vector.

```
my_vector <- c(1, 2, '3', '4', 5)
my_vector
```

```
## [1] "1" "2" "3" "4" "5"
```

Answer

In R, a vector is a sequence of data elements of the same basic type. So it automatically converts the numbers to character String and displays enclosed in double quotes.

Data import/export (3 points)

1. Download the rail_trail.txt file from Canvas (in the Midterm Exam section here) and successfully import it into R. Prove that it was imported successfully by including your import code and taking a glimpse of the result.

Answer

```
# Read from rail_trail.txt file
rail_trail <- read_csv("C:/view/opt/apps/git/R/compscix-415-2-assignments/rail_trail.txt")

## Parsed with column specification:
## cols(
##   `hightemp|lowtemp|avgtemp|spring|summer|fall|cloudcover|precip|volume|weekday` = col_character()
## )

# glimpse rail_trail
glimpse(rail_trail)

## Observations: 90
## Variables: 1
## $ `hightemp|lowtemp|avgtemp|spring|summer|fall|cloudcover|precip|volume|weekday` <chr> ...
```

2. Export the file into an R-specific format and name it “rail_trail.rds”. Make sure you define the path correctly so that you know where it gets saved. Then reload the file. Include your export and import code and take another glimpse.

Answer

```
# Read from rail_trail.txt file
rail_trail <- read_csv("C:/view/opt/apps/git/R/compscix-415-2-assignments/rail_trail.txt")

## Parsed with column specification:
## cols(
##   `hightemp|lowtemp|avgtemp|spring|summer|fall|cloudcover|precip|volume|weekday` = col_character()
## )

# glimpse rail_trail
glimpse(rail_trail)

## Observations: 90
## Variables: 1
## $ `hightemp|lowtemp|avgtemp|spring|summer|fall|cloudcover|precip|volume|weekday` <chr> ...

# Write to rail_trail.rds
saveRDS(rail_trail, "rail_trail.rds")
```

```
# load rail_trail.rds
rail_trail2 = readRDS("C:/view/opt/apps/git/R/compscix-415-2-assignments/rail_trail.rds")

# glimpse rail_trail
glimpse(rail_trail2)

## Observations: 90
## Variables: 1
## $ `hightemp|lowtemp|avgtemp|spring|summer|fall|cloudcover|precip|volume|weekday` <chr> ...
```

Visualization (6 points)

1. Critique this graphic: give only three examples of what is wrong with this graphic. Be concise.

Note: Please refer to the below link for the above mentioned graphic.

https://github.com/sanatanonline/compscix-415-2-assignments/blob/master/compscix4152_hw_5.pdf

Answer

This graphic has multiple issues which creates wrong impressions of the data visualization. The major three wrong representations are:

- This is not a standard statistical chart/plot representation (is it a bubble chart?). The numeric values definitely does not match the size of the images.
- What are those numbers represents? Are they percentage/total number of responders or what...not clear.
- What are those colors represent? No legend. Not clear.


2. Reproduce this graphic using the diamonds data set.

Note: Please refer to the below link for the above mentioned graphic.

https://github.com/sanatanonline/compscix-415-2-assignments/blob/master/compscix4152_hw_5.pdf

Answer

```
ggplot(data = diamonds, aes(x = cut, y = carat, fill = color)) +
  geom_boxplot() +
  coord_flip() +
  labs(x="CUT OF THE DIAMOND", y="CARAT OF THE DIAMOND")
```



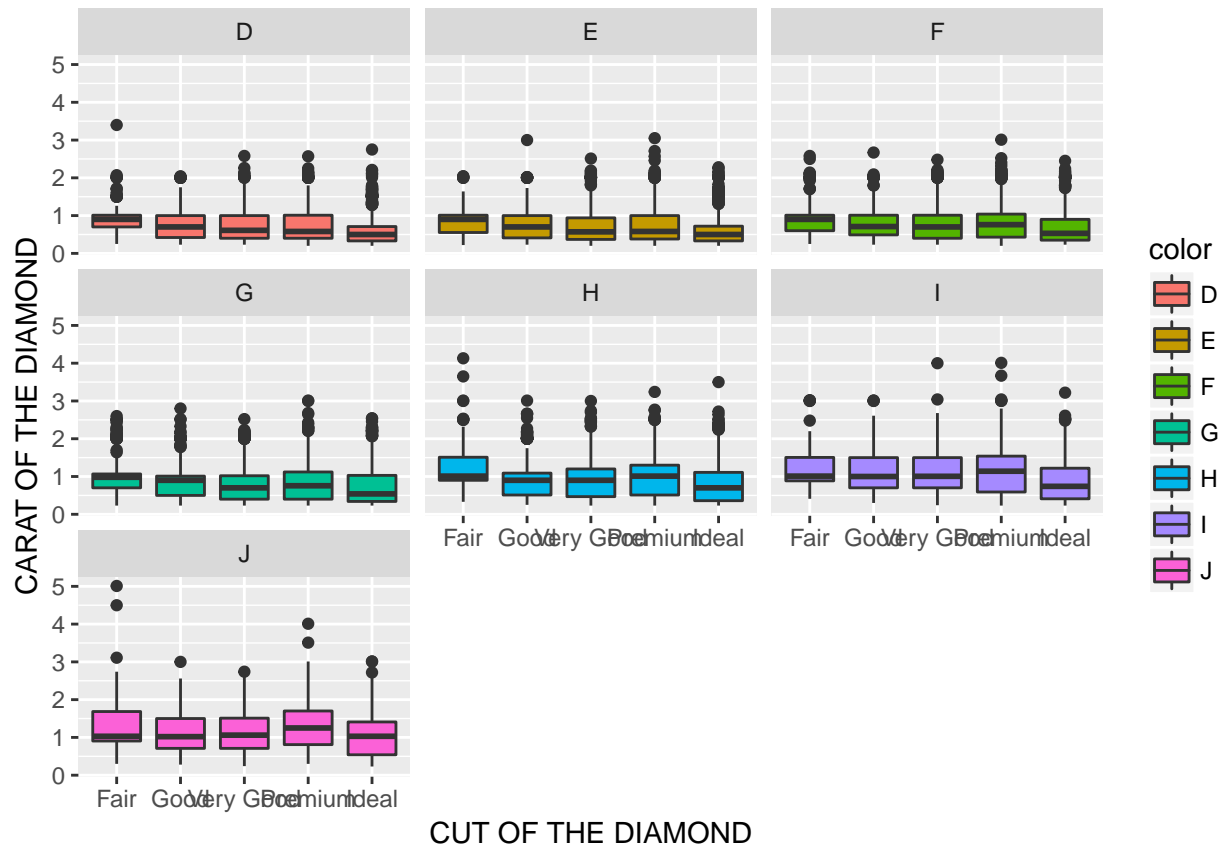
I think the box plot will be like above. But I am not sure how to get them overlapped.

3. The previous graphic is not very useful. We can make it much more useful by changing one thing about it. Make the change and plot it again.

Answer

The graphic is not very useful because the box plots are overlapped. We can make it useful by making the below change:

```
ggplot(data = diamonds, aes(x = cut, y = carat, fill = color)) +
  geom_boxplot() +
  facet_wrap(~color) +
  labs(x="CUT OF THE DIAMOND", y="CARAT OF THE DIAMOND")
```



Data munging and wrangling (6 points)

1. Is this data “tidy”? If yes, leave it alone and go to the next problem. If no, make it tidy. Note: this data set is called table2 and is available in the tidyverse package. It should be ready for you to use after you’ve loaded the tidyverse package.

```
table2
```

```
## # A tibble: 12 x 4
##   country    year type      count
##   <chr>      <int> <chr>      <int>
## 1 Afghanistan 1999 cases        745
## 2 Afghanistan 1999 population 19987071
## 3 Afghanistan 2000 cases        2666
## 4 Afghanistan 2000 population 20595360
## 5 Brazil      1999 cases        37737
## 6 Brazil      1999 population 172006362
## 7 Brazil      2000 cases        80488
## 8 Brazil      2000 population 174504898
## 9 China       1999 cases        212258
## 10 China      1999 population 1272915272
## 11 China      2000 cases        213766
## 12 China      2000 population 1280428583
```

Answer

This is not a tidy data. This dataset intermingles the values of population and cases in the same columns. As a result, we would need to untangle the values whenever we want to work with each variable separately.

The *key* column contains only keys (and not just because the column is labelled *key*). Conveniently, the *value* column contains the values associated with those keys.

We can use the *spread()* function to tidy this layout. So the tidy form of the dataset would be like below:

```
spread(table2, type, count)
```

```
## # A tibble: 6 x 4
##   country      year cases population
##   <chr>      <int> <int>      <int>
## 1 Afghanistan 1999     745   19987071
## 2 Afghanistan 2000    2666  20595360
## 3 Brazil      1999   37737  172006362
## 4 Brazil      2000   80488  174504898
## 5 China       1999 212258 1272915272
## 6 China       2000 213766 1280428583
```

2. Create a new column in the diamonds data set called *price_per_carat* that shows the price of each diamond per carat (hint: divide). Only show me the code, not the output.

We can do this using the code below:

```
mutate(diamonds, price_per_carat <- price / carat)
```

3. For each cut of diamond in the diamonds data set, how many diamonds, and what proportion, have a price > 10000 and a carat < 1.5? There are several ways to get to an answer, but your solution must use the data wrangling verbs from the tidyverse in order to get credit.

- Do the results make sense? Why?
- Do we need to be wary of any of these numbers? Why?

Answer

```
filter(diamonds, price > 10000, carat < 1.5)
```

```
## # A tibble: 834 x 10
##   carat cut      color clarity depth table price      x      y      z
##   <dbl> <ord>    <ord> <ord>    <dbl> <dbl> <int> <dbl> <dbl> <dbl>
## 1  1.03 Ideal      E     VVS2    60.6  59.0 10003  6.50  6.53  3.95
## 2  1.23 Very Good G     VVS2    60.6  55.0 10004  6.93  7.02  4.23
## 3  1.25 Ideal      F     VS2     61.6  55.0 10006  6.93  6.96  4.28
## 4  1.21 Very Good F     VS1     62.3  58.0 10009  6.76  6.85  4.24
## 5  1.01 Fair       D     SI2     64.6  58.0 10011  6.25  6.20  4.02
## 6  1.05 Ideal      F     VVS2    60.5  55.0 10011  6.67  6.58  4.01
## 7  1.35 Premium    G     VS1     62.1  59.0 10012  7.06  7.02  4.37
## 8  1.13 Ideal      F     VS1     60.9  57.0 10016  6.73  6.76  4.11
## 9  1.21 Premium    F     VS1     62.6  59.0 10018  6.81  6.76  4.25
## 10 1.01 Very Good F     VVS1    62.9  57.0 10019  6.35  6.41  4.01
## # ... with 824 more rows
```

```
diamonds2 <- diamonds %>%
  group_by(cut) %>%
  summarise(prop = sum(price > 10000, carat < 1.5) / n()) %>%
  arrange(cut)

print(tbl_df(diamonds2))
```



```
## # A tibble: 5 x 2
##   cut      prop
##   <ord>    <dbl>
## 1 Fair      0.889
## 2 Good      0.946
## 3 Very Good 0.979
## 4 Premium   0.969
## 5 Ideal     1.01
```

This table shows almost the whole proportion of each cut have price more than 10000 when carat < 1.5? Confusing.

EDA (6 points)

Take a look at the *txhousing* data set that is included with the *ggplot2* package and answer these questions:

1. During what time period is this data from?
2. How many cities are represented?
3. Which city, month and year had the highest number of sales?
4. What kind of relationship do you think exists between the number of listings and the number of sales? Check your assumption and show your work.
5. What proportion of sales is missing for each city?
6. Looking at only the cities and months with greater than 500 sales:
 - Are the distributions of the median sales price (column name median), when grouped by city, different? The same? Show your work.
 - Any cities that stand out that you'd want to investigate further?
 - Why might we want to filter out all cities and months with sales less than 500?

Answer

To do the EDA on *txhousing* data, first we take a quick look at the dataset. We can use `?txhousing` for help to understand the variables.

```
txhousing

## # A tibble: 8,602 x 9
##   city      year month sales  volume median listings inventory date
##   <chr>    <int> <int> <dbl>    <dbl> <dbl>    <dbl>    <dbl> <dbl>
## 1 Abilene  2000     1  72.0  5380000 71400     701     6.30 2000
## 2 Abilene  2000     2  98.0  6505000 58700     746     6.60 2000
## 3 Abilene  2000     3 130    9285000 58100     784     6.80 2000
## 4 Abilene  2000     4  98.0  9730000 68600     785     6.90 2000
## 5 Abilene  2000     5 141   10590000 67300     794     6.80 2000
## 6 Abilene  2000     6 156   13910000 66900     780     6.60 2000
## 7 Abilene  2000     7 152   12635000 73500     742     6.20 2000
## 8 Abilene  2000     8 131   10710000 75000     765     6.40 2001
## 9 Abilene  2000     9 104    7615000 64500     771     6.50 2001
## 10 Abilene 2000    10 101    7040000 59300     764     6.60 2001
## # ... with 8,592 more rows
```

From the above result, we see its a dataset of 9 variables with 8602 observations. Now we will do the analysis to answer the above questions.

1. During what time period is this data from?

```
arrange(txhousing, year, month)
```

```
## # A tibble: 8,602 x 9
##   city      year month  sales  volume median listings inventory date
##   <chr>    <int> <int> <dbl>   <dbl>   <dbl>   <dbl>   <dbl> <dbl>
## 1 Abilene    2000     1   72.0  5.38e6  71400     701     6.30  2000
## 2 Amarillo   2000     1   102   8.86e6  80000     972     5.30  2000
## 3 Arlington  2000     1   241   2.62e7  94000    1417     3.70  2000
## 4 Austin     2000     1  1025   1.73e8 133700    3084     2.00  2000
## 5 Bay Area   2000     1   244   2.93e7 100700    1766     4.30  2000
## 6 Beaumont   2000     1    97.0  1.01e7  82100     876     6.10  2000
## 7 Brazoria Co~ 2000     1    55.0  5.24e6  74400     512     5.90  2000
## 8 Brownsville 2000     1    NA    NA        NA     400     9.10  2000
## 9 Bryan-Colle~ 2000     1    61.0  5.61e6  77900     498     4.20  2000
## 10 Collin Coun~ 2000     1   464   9.48e7 158700    2844     4.00  2000
## # ... with 8,592 more rows
```

```
arrange(txhousing, desc(year), desc(month))
```

```
## # A tibble: 8,602 x 9
##   city      year month  sales  volume median listings inventory date
##   <chr>    <int> <int> <dbl>   <dbl>   <dbl>   <dbl>   <dbl> <dbl>
## 1 Abilene    2015     7    268  4.58e7 148700     986     5.00  2016
## 2 Amarillo   2015     7    354  6.23e7 149700    1247     4.50  2016
## 3 Arlington  2015     7    605  1.25e8 178900     752     1.70  2016
## 4 Austin     2015     7   3466  1.15e9 264600    7913     3.00  2016
## 5 Bay Area   2015     7    849  1.97e8 200800    2144     3.20  2016
## 6 Beaumont   2015     7    318  5.29e7 139300    1561     6.40  2016
## 7 Brazoria Co~ 2015     7    NA    NA        NA     NA     NA     2016
## 8 Brownsville 2015     7    NA    NA        NA     NA     NA     2016
## 9 Bryan-Colle~ 2015     7    414  9.04e7 190700     894     3.30  2016
## 10 Collin Coun~ 2015     7   1861  6.14e8 292600    2809     2.10  2016
## # ... with 8,592 more rows
```

From the above results, we see that the data is collected monthly from Jan 2000 to July 2015

2. How many cities are represented?

```
count(unique(txhousing[,1]))
```

```
## # A tibble: 1 x 1
##       n
##   <int>
## 1    46
```

There are 46 cities represented in *txhousing* dataset.

3. Which city, month and year had the highest number of sales?

In this dataset *sales* variable represents the **number of sales**. So we arrange the dataset in descending order by sales.

```
arrange(txhousing, desc(sales))
```

```
## # A tibble: 8,602 x 9
```


```
##   city    year month sales    volume median listings inventory date
##   <chr>  <int> <int> <dbl>    <dbl>   <dbl>   <dbl>    <dbl> <dbl>
## 1 Houston  2015     7  8945 2568156780 217600   23875     3.40 2016
## 2 Houston  2006     6  8628 1795898108 155200   36281     5.60 2006
## 3 Houston  2013     7  8468 2168720825 187800   21497     3.30 2014
## 4 Houston  2015     6  8449 2490238594 222400   22311     3.20 2015
## 5 Houston  2013     5  8439 2121508529 186100   20526     3.30 2013
## 6 Houston  2014     6  8391 2342443127 211200   19725     2.90 2014
## 7 Houston  2014     7  8391 2278932511 199700   20214     3.00 2014
## 8 Houston  2014     8  8167 2195184825 202400   20007     2.90 2015
## 9 Houston  2013     8  8155 2083377894 186700   21366     3.30 2014
## 10 Houston 2006     5  8040 1602621368 151200   35398     5.50 2006
## # ... with 8,592 more rows
```

In the above result we see, Houston had the highest number of sales 8945 in July (month 7), 2015. This city had maximum volume (total value of sales) too.

4. What kind of relationship do you think exists between the number of listings and the number of sales? Check your assumption and show your work.

```
ggplot(data = txhousing) +
  geom_point(mapping = aes(x = listings, y = sales))
```

```
## Warning: Removed 1426 rows containing missing values (geom_point).
```




In the above scatter plot we see that there is a relationship between the number of listings and the number of sales. When there is more number of listings, number of sales increases. The below plot confirms the trend.

```
ggplot(data = txhousing) +
  geom_smooth(mapping = aes(x = listings, y = sales))
```

```
## `geom_smooth()` using method = 'gam'
```

```
## Warning: Removed 1426 rows containing non-finite values (stat_smooth).
```



5. What proportion of sales is missing for each city?

We can find out the proportion of sales is missing for each city using following code.

```
missing_sales_prop_per_city <- txhousing %>%
  group_by(city) %>%
  summarise(prop = sum(is.na(sales)) / n()) %>%
  arrange(city)

# we have to print 46 rows for 46 cities
print(tbl_df(missing_sales_prop_per_city), n=46)
```

```
## # A tibble: 46 x 2
##   city                prop
##   <chr>              <dbl>
## 1 Abilene            0
## 2 Amarillo           0
## 3 Arlington          0
## 4 Austin             0
## 5 Bay Area           0
## 6 Beaumont           0
```

```
## 7 Brazoria County      0.0749
## 8 Brownsville          0.0107
## 9 Bryan-College Station 0
## 10 Collin County       0
## 11 Corpus Christi      0.00535
## 12 Dallas              0
## 13 Denton County       0
## 14 El Paso             0
## 15 Fort Bend           0
## 16 Fort Worth          0
## 17 Galveston           0.00535
## 18 Garland             0
## 19 Harlingen           0.134
## 20 Houston             0
## 21 Irving              0
## 22 Kerrville           0.556
## 23 Killeen-Fort Hood   0.00535
## 24 Laredo              0.193
## 25 Longview-Marshall   0.0642
## 26 Lubbock             0.00535
## 27 Lufkin              0
## 28 McAllen             0.0107
## 29 Midland             0.401
## 30 Montgomery County   0
## 31 Nacogdoches         0.0588
## 32 NE Tarrant County   0
## 33 Odessa              0.385
## 34 Paris               0
## 35 Port Arthur         0.0107
## 36 San Angelo          0
## 37 San Antonio         0
## 38 San Marcos          0.246
## 39 Sherman-Denison     0
## 40 South Padre Island   0.620
## 41 Temple-Belton       0.0588
## 42 Texarkana           0.0909
## 43 Tyler               0
## 44 Victoria            0
## 45 Waco                 0.102
## 46 Wichita Falls       0
```

6. Looking at only the cities and months with greater than 500 sales:

- Are the distributions of the median sales price (column name median), when grouped by city, different? The same? Show your work.
- Any cities that stand out that you'd want to investigate further?
- Why might we want to filter out all cities and months with sales less than 500?

First filter out the months with less than or equal to 500 as we are going to look at only the cities and months with greater than 500 sales.


```
newtxhousing <- filter(txhousing, sales > 500)

newtxhousing
```

```
## # A tibble: 1,883 x 9
##   city      year month sales  volume median listings inventory date
##   <chr>    <int> <int> <dbl>    <dbl>    <dbl>    <dbl>    <dbl> <dbl>
## 1 Arlington 2000     8  507 60875199 103400    1417    3.50 2001
## 2 Arlington 2001     5  536 69878959 114400    1592    3.70 2001
## 3 Arlington 2001     6  534 67744182 108500    1627    3.80 2001
## 4 Arlington 2001     8  505 65080743 113600    1616    3.70 2002
## 5 Arlington 2002     5  503 67240236 116100    1741    3.90 2002
## 6 Arlington 2002     7  509 66954143 119100    1925    4.40 2002
## 7 Arlington 2003     5  502 67131982 118000    2544    5.90 2003
## 8 Arlington 2003     7  524 73194692 123500    2799    6.50 2004
## 9 Arlington 2003     8  531 72397143 123900    2801    6.40 2004
## 10 Arlington 2004     5  527 72401436 118300    2922    6.50 2004
## # ... with 1,873 more rows
```

And we see there are such 1883 records.

```
ggplot(newtxhousing, aes(date, median)) +
  geom_line(aes(group = city), alpha = 1/2) +
  geom_line(stat = "summary", fun.y = "mean", colour = "red")
```



Looking at the above plot, we can see the some of the cities (I guess big cities like Houston) has more sales than other cities.

```
#OutVals = boxplot(newtxhousing$sales)$out
#which(newtxhousing %in% OutVals)

OutVals <- tibble(boxplot(newtxhousing$median, plot=FALSE)$out)
```

```
OutVals[1]
```

```
## # A tibble: 56 x 1
##   `boxplot(newtxhousing$median, plot = FALSE)$out`
##                                     <dbl>
## 1                                     248900
## 2                                     246900
## 3                                     245700
## 4                                     245300
## 5                                     253900
## 6                                     270300
## 7                                     271200
## 8                                     270200
## 9                                     264600
## 10                                    252600
## # ... with 46 more rows
```

```
outliers_cities <- subset(newtxhousing, newtxhousing$median %in% OutVals)
```

```
outliers_cities
```

```
## # A tibble: 0 x 9
## # ... with 9 variables: city <chr>, year <int>, month <int>, sales <dbl>,
## #   volume <dbl>, median <dbl>, listings <dbl>, inventory <dbl>,
## #   date <dbl>
```

```
filter(newtxhousing, newtxhousing$median %in% OutVals[1])
```


```
## # A tibble: 0 x 9
## # ... with 9 variables: city <chr>, year <int>, month <int>, sales <dbl>,
## #   volume <dbl>, median <dbl>, listings <dbl>, inventory <dbl>,
## #   date <dbl>
```

First lets see the records having sales less than 500.

```
txhousing2 <- filter(txhousing, sales < 500)
```

```
ggplot(data = txhousing2) +
  geom_smooth(mapping = aes(x = txhousing2$year, y = txhousing2$sales))
```

```
## `geom_smooth()` using method = 'gam'
```



There is no much change for the records having sales less than 500.

Git and Github (1.5 points)

To demonstrate your use of git and Github, at the top of your document put a hyperlink to your Github repository.

Answer

All the work is pushed to Github repository. The repository URL is below.

<https://github.com/sanatanonline/compscix-415-2-assignments>

End of Homework 5/Midterm